

Self-Driving Cars

Lecture 2 - Imitation Learning

Prof. Dr.-Ing. Andreas Geiger

Autonomous Vision Group
MPI-IS / University of Tübingen

October 25, 2018



University of Tübingen
MPI for Intelligent Systems
Autonomous Vision Group

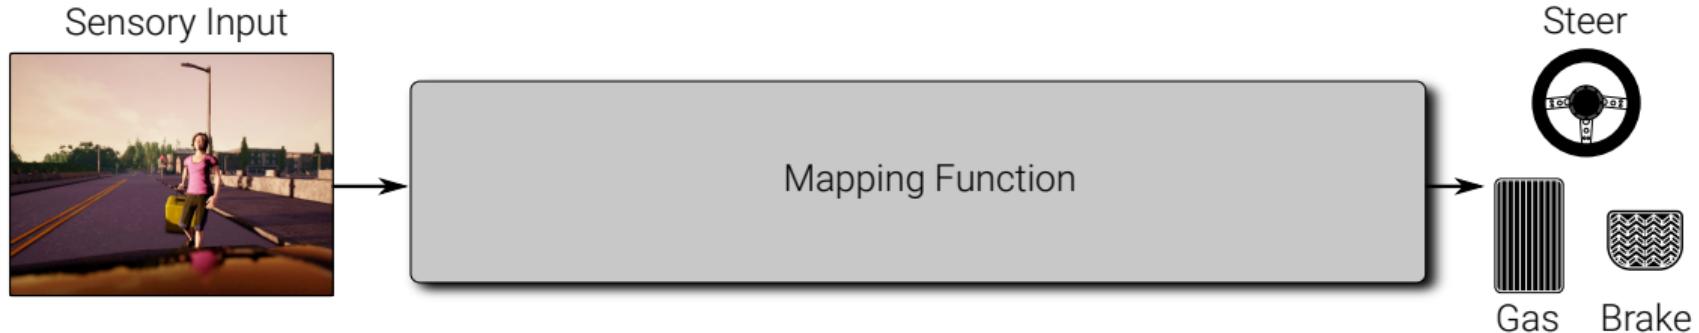


Agenda

Date	Lecture (Thursday)	Date	Exercise (Friday)
18.10.	01 - Introduction to Self-Driving Cars	19.10.	00 - Introduction Pytorch & OpenAI Gym
25.10.	02 - DNNs, ConvNets, Imitation Learning	26.10.	01 - Intro: Imitation Learning
1.11.	none (Allerheiligen)	2.11.	
8.11.	03 - Direct Perception	9.11.	01 - Q&A
15.11.	none (CVPR Deadline)	16.11.	
22.11.	04 - Reinforcement Learning	23.11.	01 - Discussion & 02 - Intro: Reinforcement Learning
29.11.	05 - Vehicle Dynamics & Control	30.11.	
6.12.	06 - Localization & Visual Odometry	7.12.	02 - Q&A
13.12.	07 - Simultaneous Localization and Mapping (J. Stückler)	14.12.	
20.12.	08 - Road and Lane Detection	21.12.	02 - Discussion & 03 - Intro: Modular Pipeline
10.1.	09 - Reconstruction and Motion Estimation	11.1.	
17.1.	10 - Object Detection & Tracking	18.1.	
24.1.	11 - Scene Understanding	25.1.	03 - Q&A
31.1.	12 - Planning	1.2.	03 - Discussion & Announcement of Winners
7.2.	13 - Winner's Presentations and Exam Q&A	8.2.	

Approaches to Self-Driving

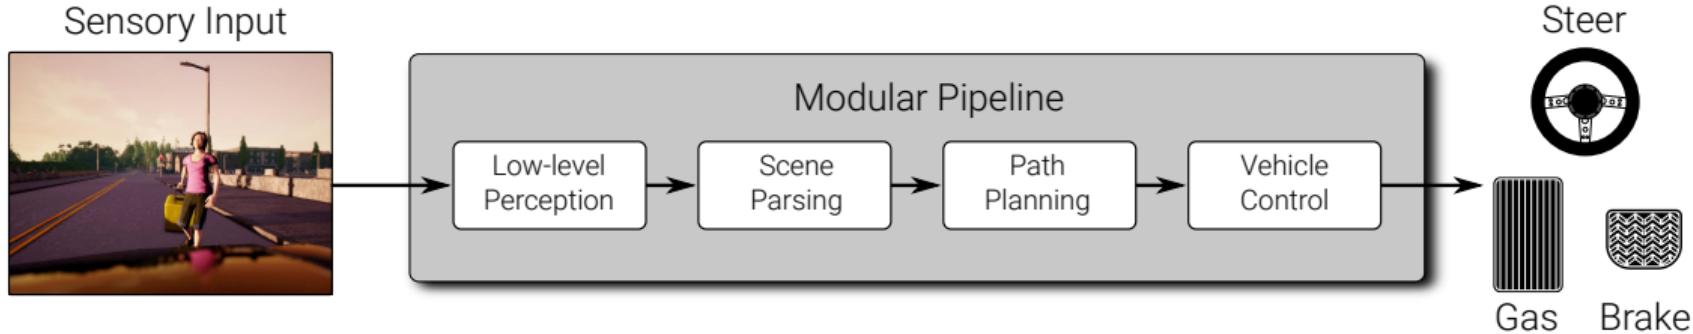
Autonomous Driving



Dominating Paradigms:

- ▶ Modular Pipelines
- ▶ End-to-End Learning (Imitation Learning, Reinforcement Learning)
- ▶ Direct Perception

Autonomous Driving: Modular Pipeline

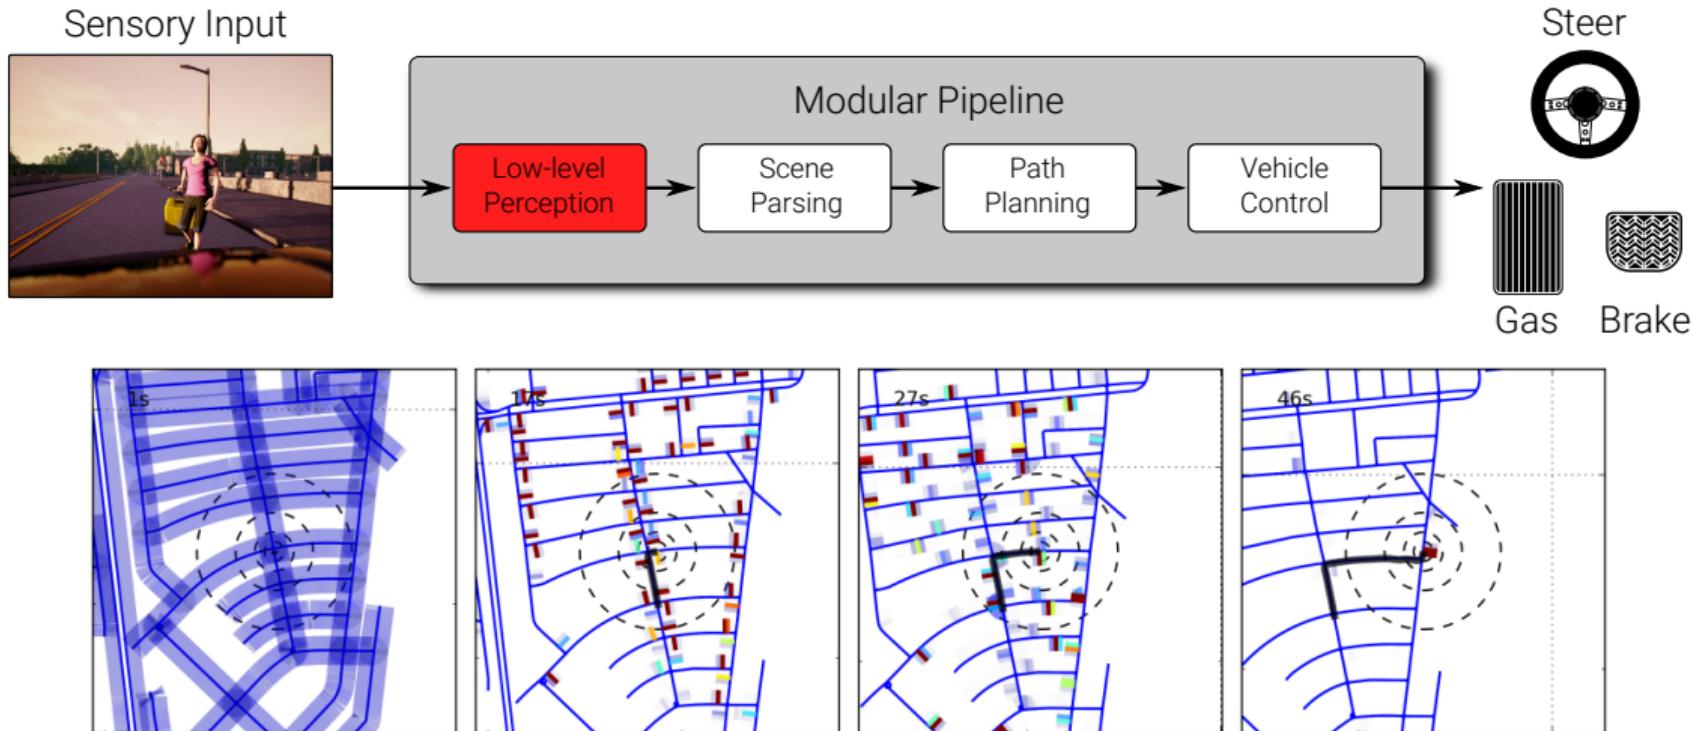


Examples:

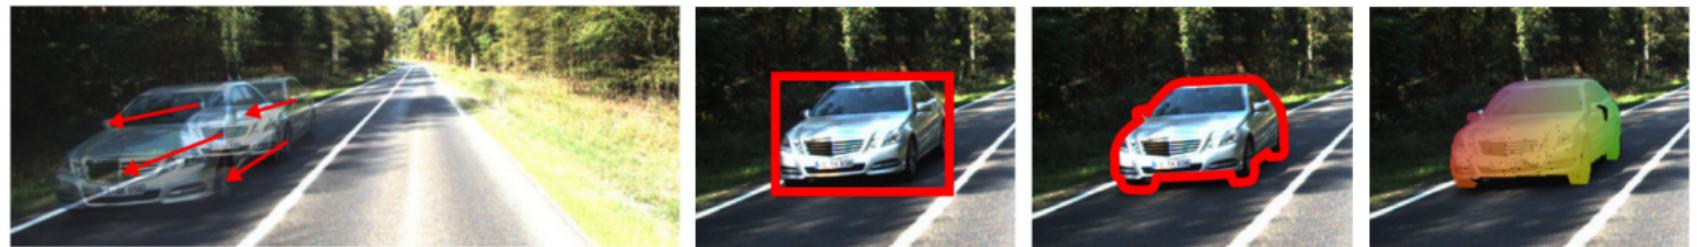
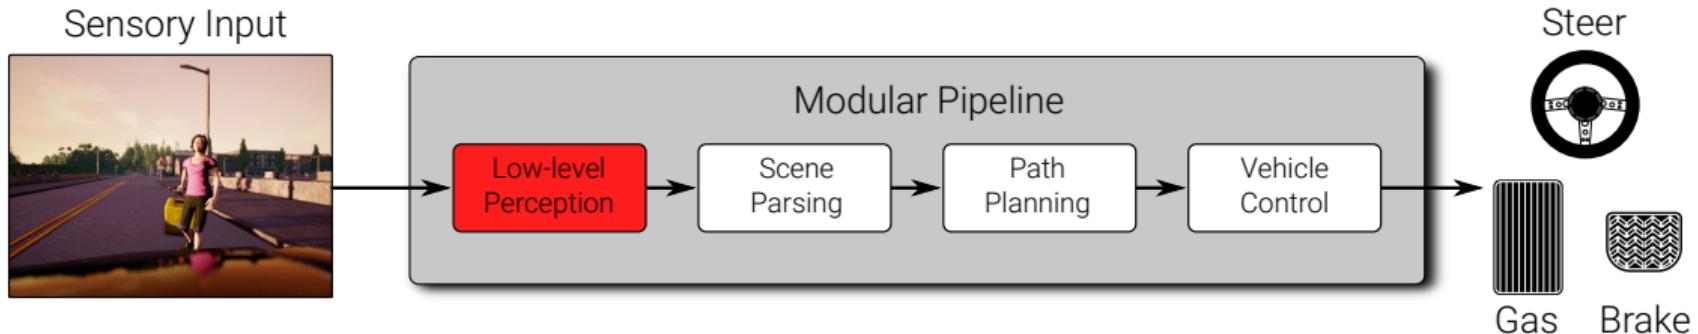
- ▶ [Montemerlo et al., JFR 2008]
- ▶ [Urmson et al., JFR 2008]
- ▶ Waymo, Uber, Tesla, Zoox, ...



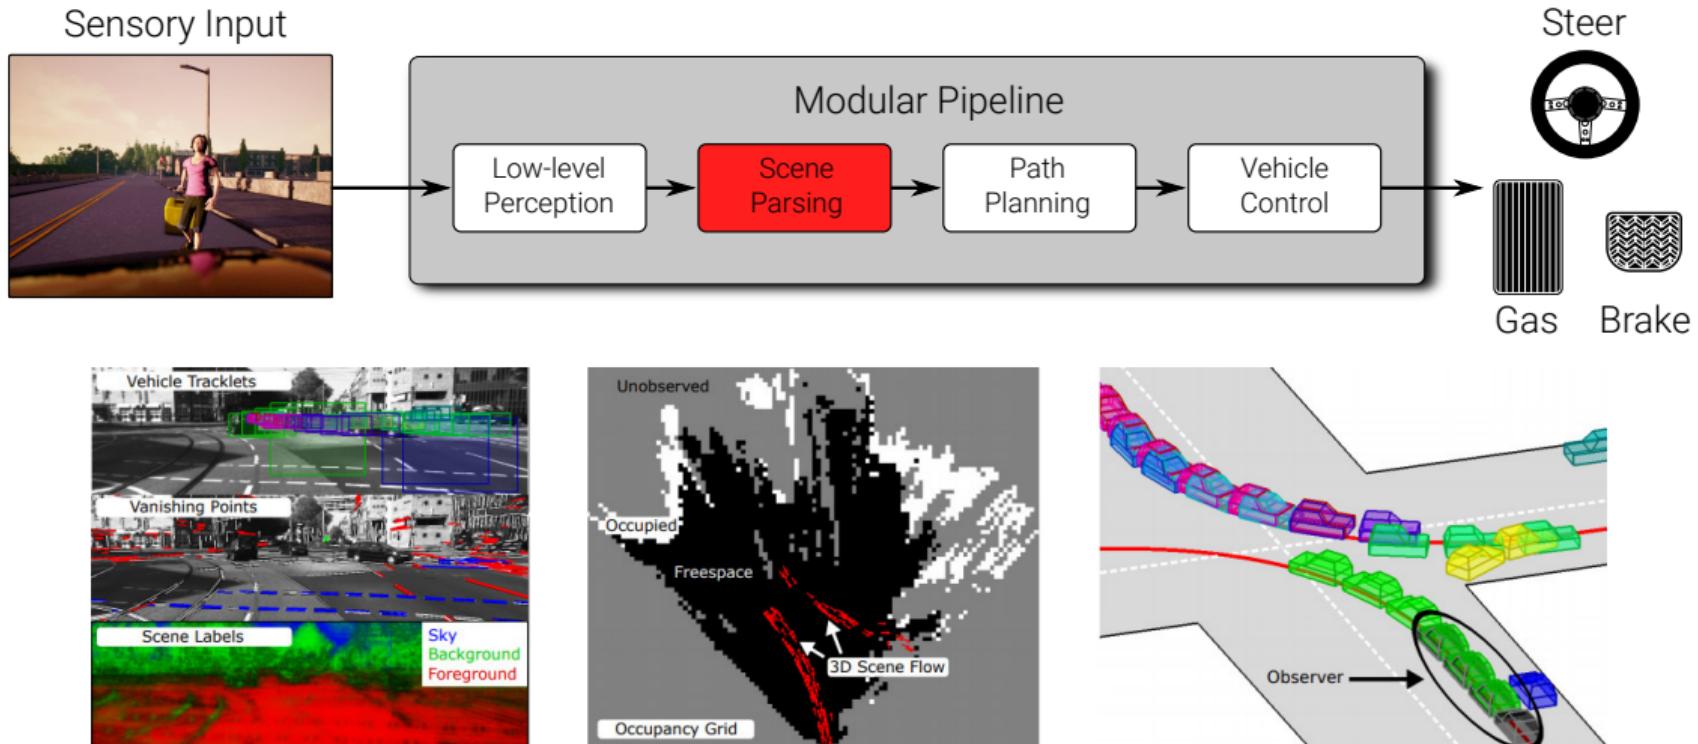
Autonomous Driving: Modular Pipeline



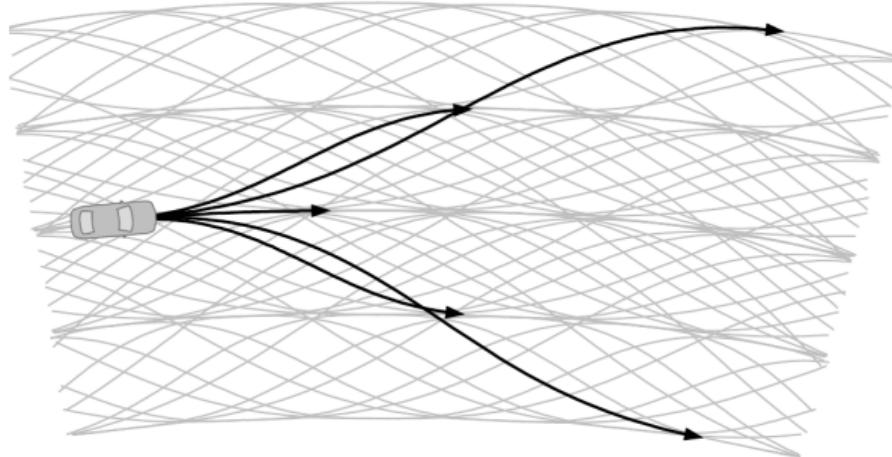
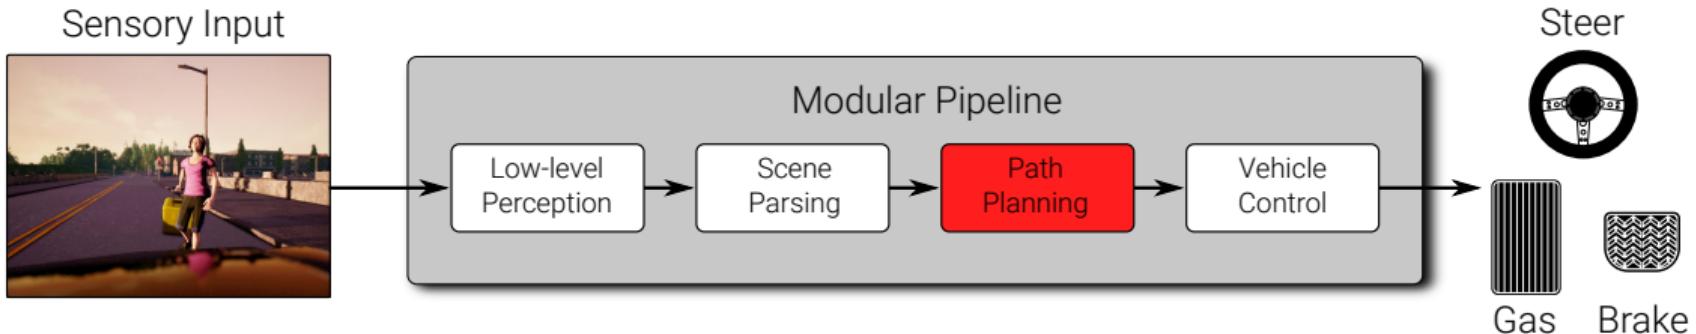
Autonomous Driving: Modular Pipeline



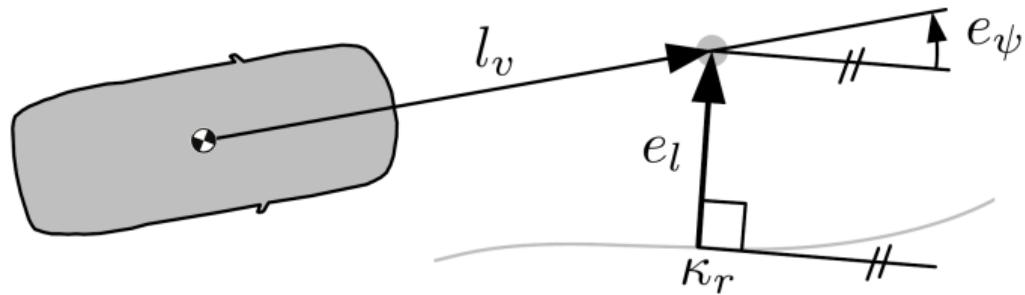
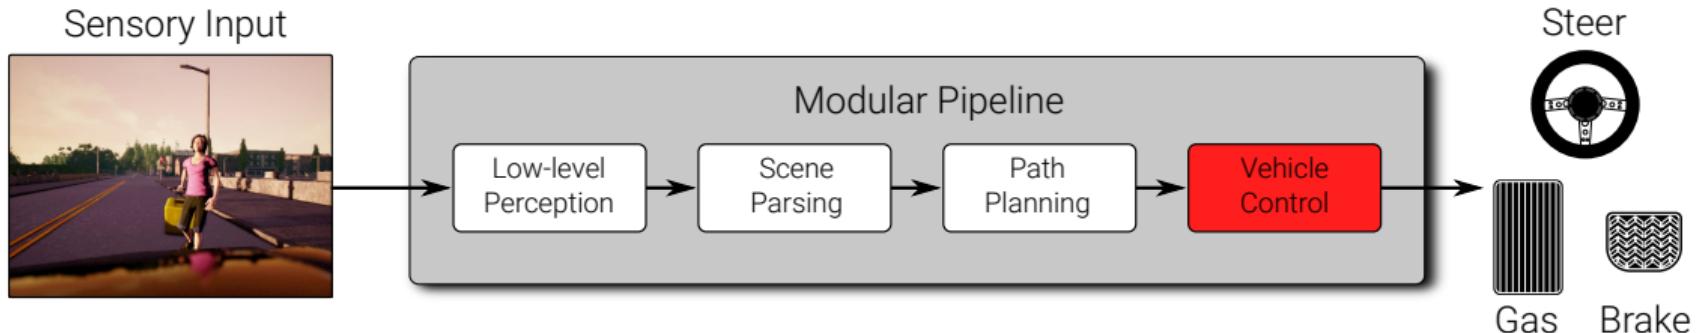
Autonomous Driving: Modular Pipeline



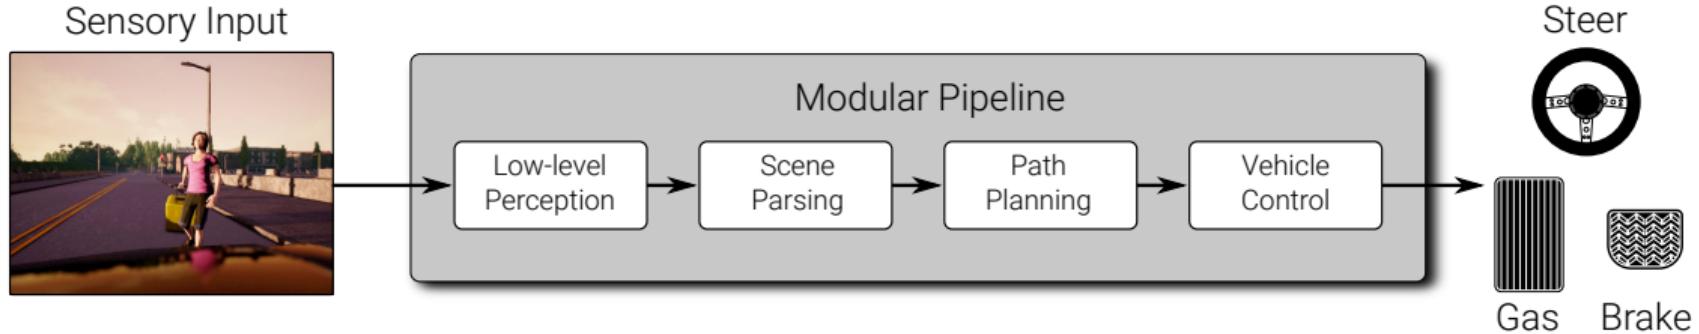
Autonomous Driving: Modular Pipeline



Autonomous Driving: Modular Pipeline



Autonomous Driving: Modular Pipeline



Pros:

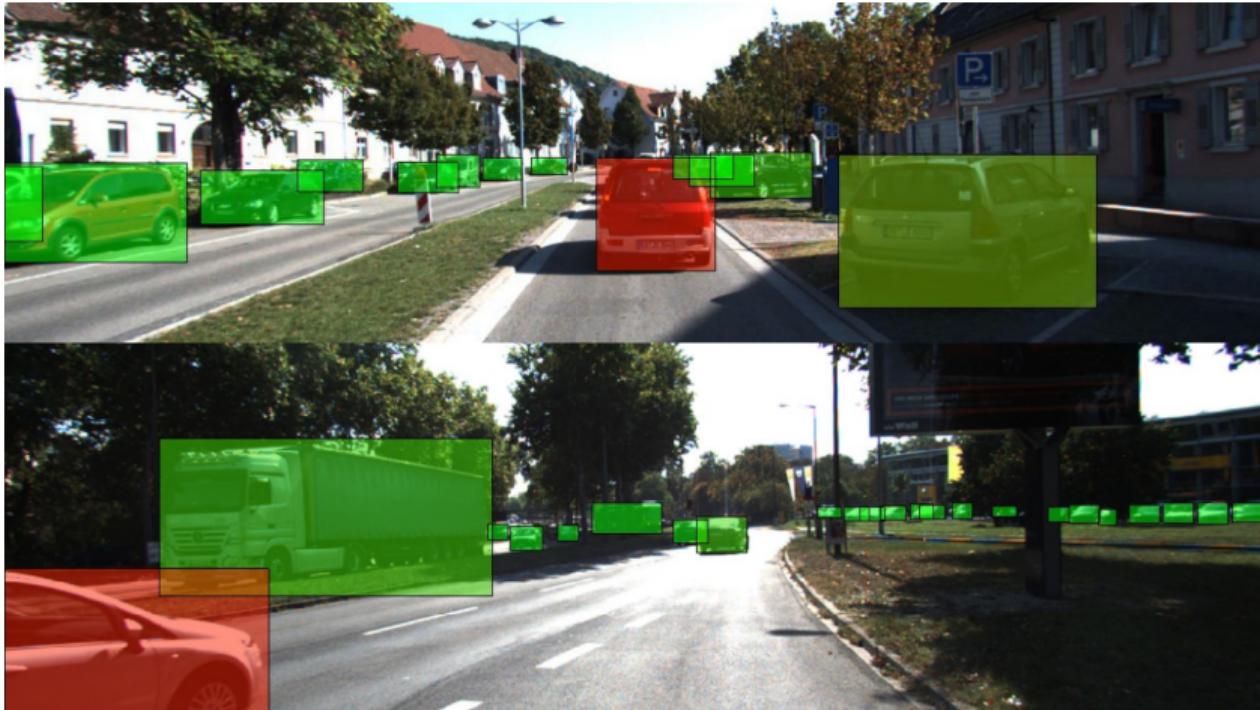
- ▶ Small components, easy to develop in parallel
- ▶ Interpretability

Cons:

- ▶ Piece-wise training (not jointly)
- ▶ Localization and planning heavily relies on HD maps

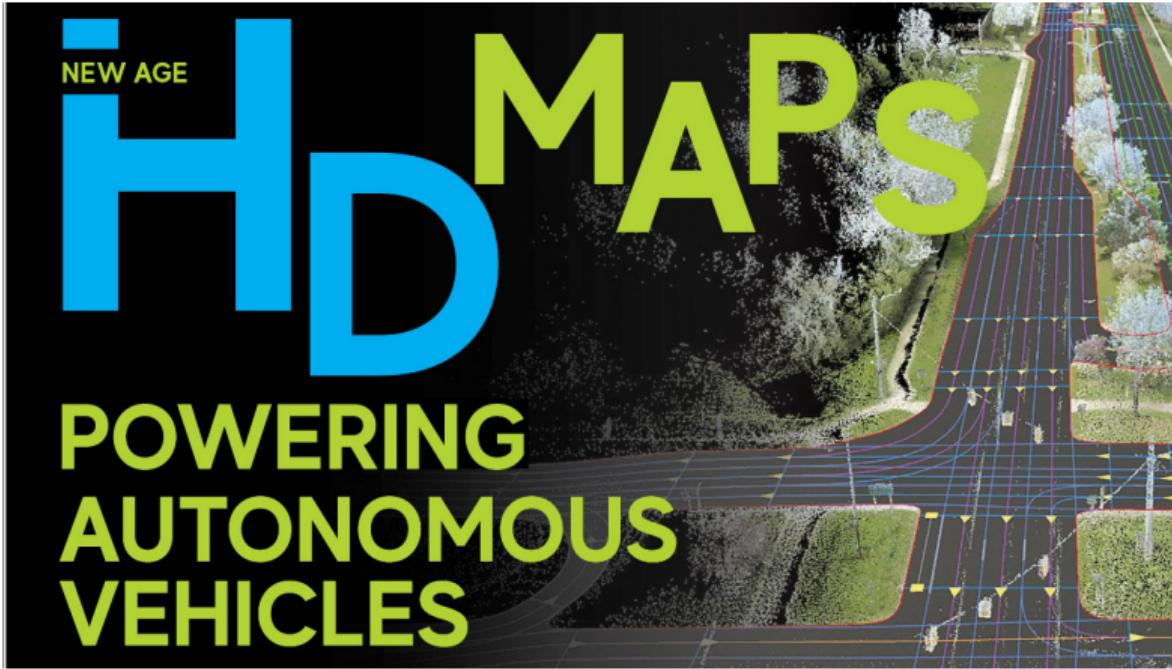
HD maps: Centimeter precision lanes, markings, traffic lights/signs, human annotated

Autonomous Driving: Modular Pipeline



- ▶ Piece-wise training difficult: not all objects are equally important!

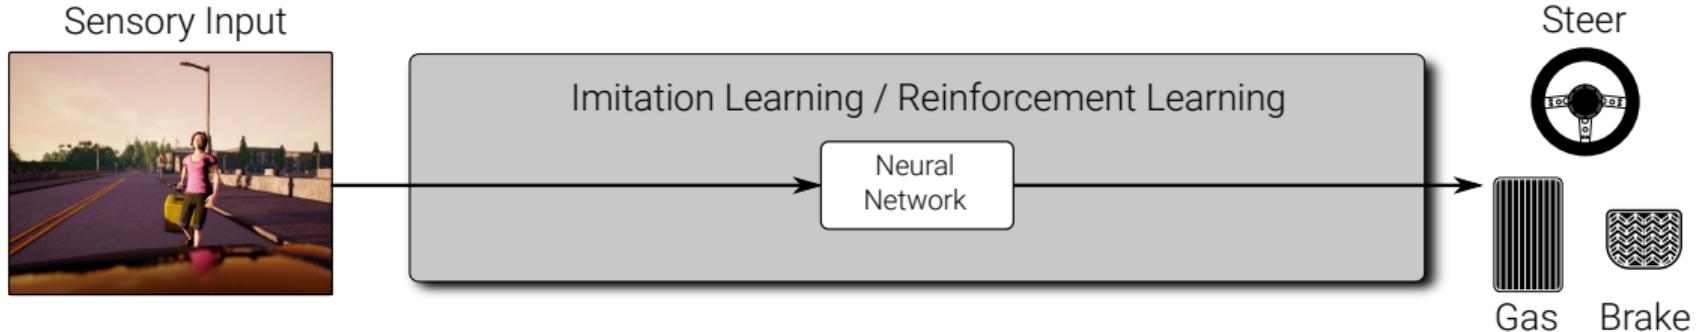
Autonomous Driving: Modular Pipeline



<https://www.geospatialworld.net/article/hd-maps-autonomous-vehicles/>

- ▶ HD Maps are expensive to create (data collection & annotation effort)

Autonomous Driving: End-to-End Learning

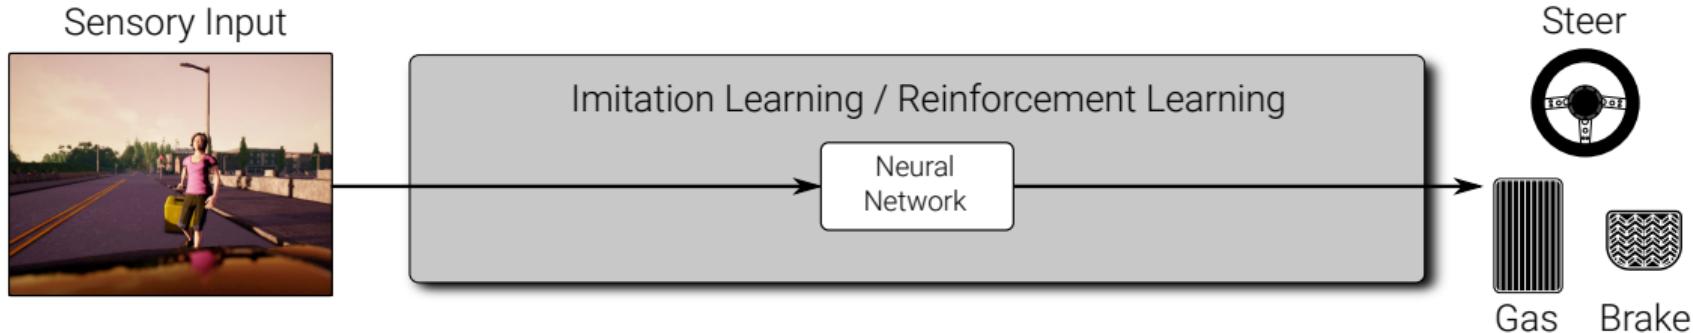


Examples:

- ▶ [Pomerleau, NIPS 1989]
- ▶ [Bojarski, Arxiv 2016]
- ▶ [Codevilla et al., ICRA 2018]



Autonomous Driving: End-to-End Learning



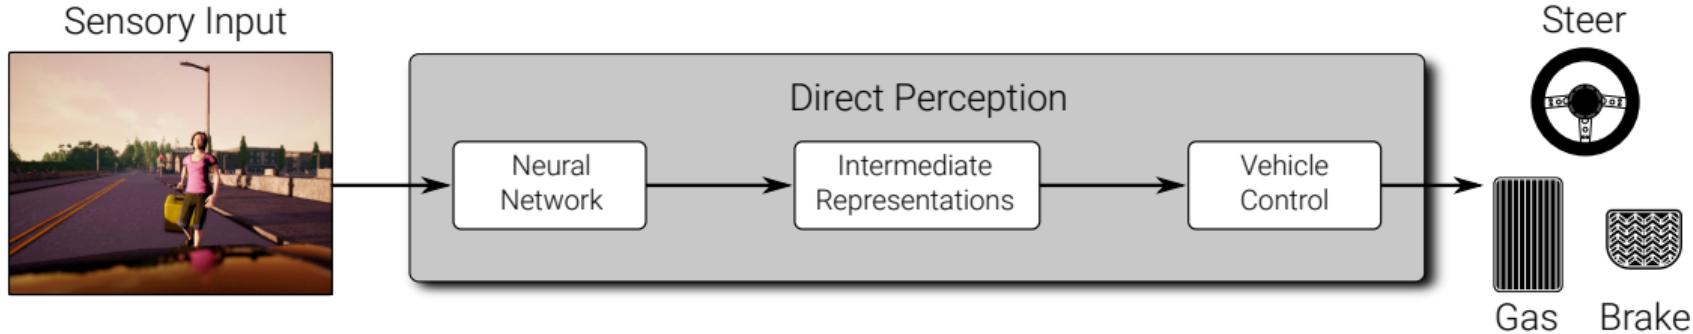
Pros:

- ▶ End-to-end training
- ▶ Cheap annotations

Cons:

- ▶ Training / Generalization
- ▶ Interpretability

Autonomous Driving: Direct Perception

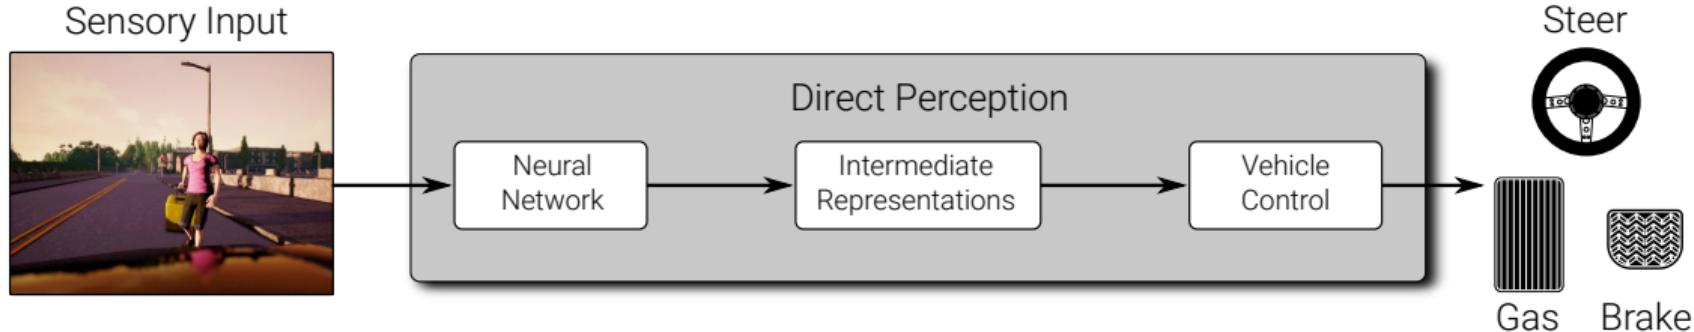


Examples:

- [Chen et al., ICCV 2015]
- [Sauer et al., CoRL 2018]



Autonomous Driving: Direct Perception



Pros:

- ▶ Compact Representation
- ▶ Interpretability

Cons:

- ▶ Control typically not learned jointly
- ▶ How to choose representations?

Topic for Today:
Imitation Learning

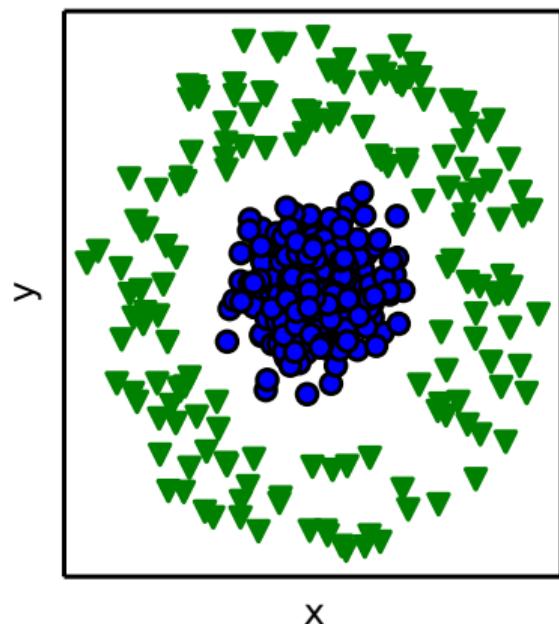
But before a little recap on ...



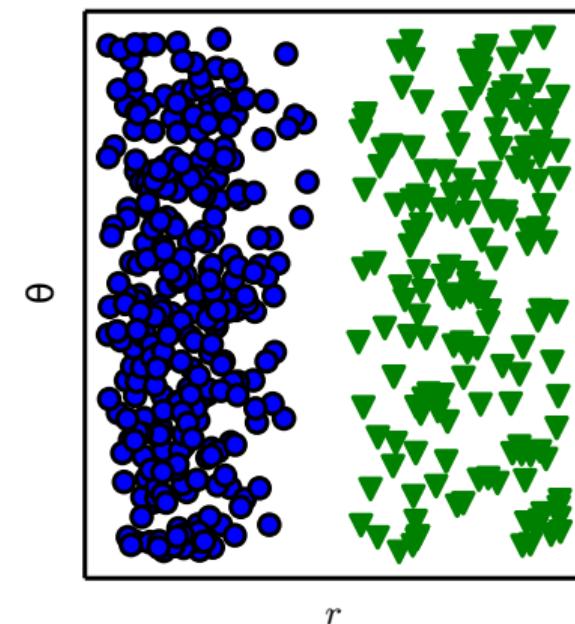
Deep Learning

Motivation: Representation Matters

Cartesian Coordinates

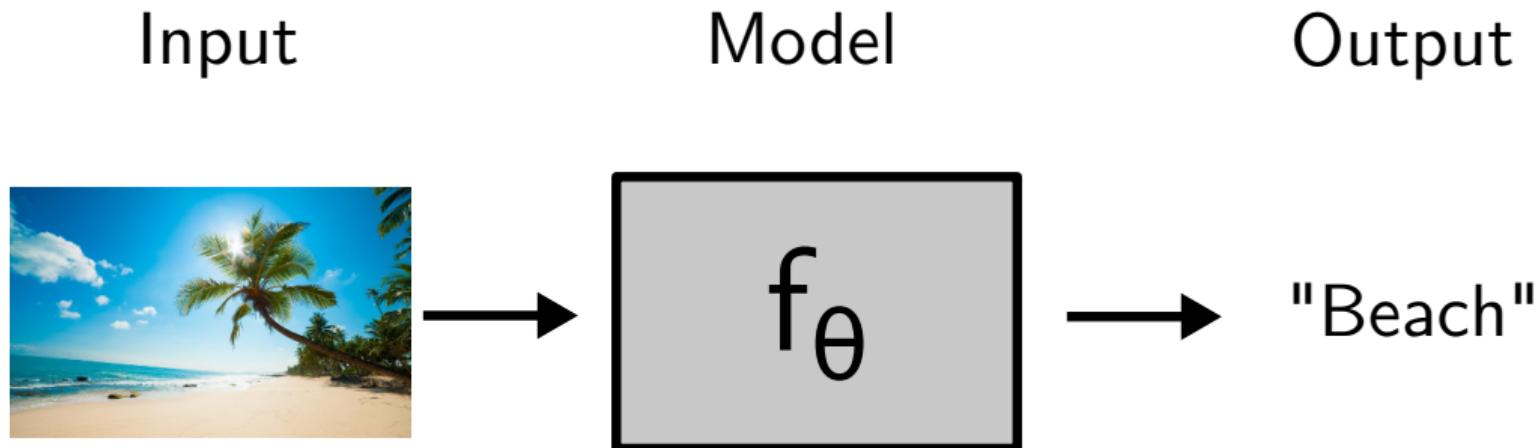


Polar Coordinates



- ▶ How to separate the blue points from the green points with a linear classifier?

Example: Multi-Class Classification



- ▶ $f_{\theta} : \mathbf{x} \in \mathbb{R}^{W \times H} \mapsto y \in \{1, \dots, L\}$
- ▶ $f_{\theta} : \mathbf{x} \in \mathbb{R}^{W \times H} \mapsto \mathbf{y} = [0, \dots, \textcolor{red}{1}, \dots, 0] \in \{0, 1\}^L$

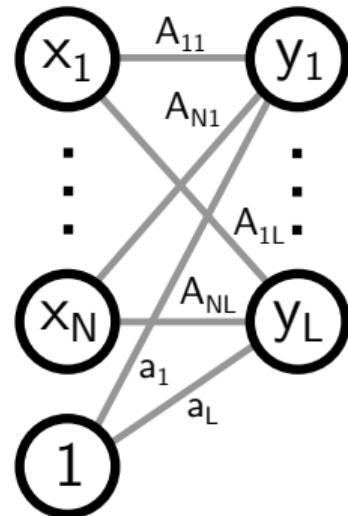
Linear Regression

- ▶ Mapping:
 \mathbf{x} = Image



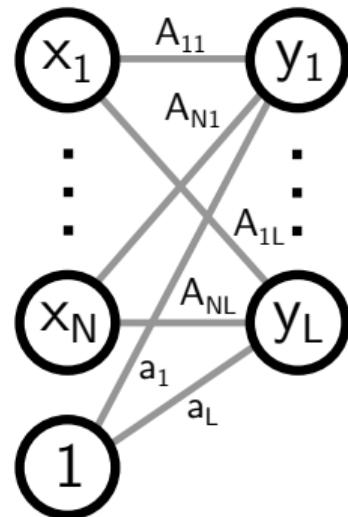
Linear Regression

- ▶ Mapping:
 \mathbf{x} = Image
 $\mathbf{y} = \mathbf{Ax} + \mathbf{a}$



Linear Regression

- ▶ Mapping:
 \mathbf{x} = Image
 $\mathbf{y} = \mathbf{Ax} + \mathbf{a}$
- ▶ Classification:
 $L^* = \operatorname{argmax}_l y_l$



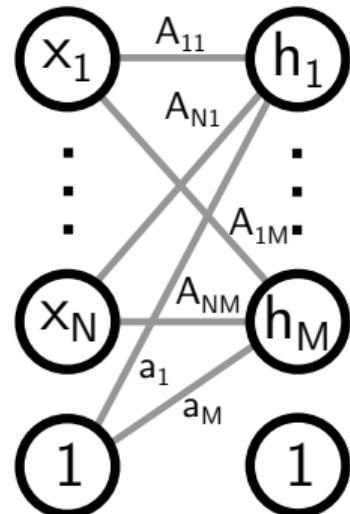
Linear Regression

- 3 Layers:
 \mathbf{x} = Image



Linear Regression

- 3 Layers:
 - \mathbf{x} = Image
 - $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$



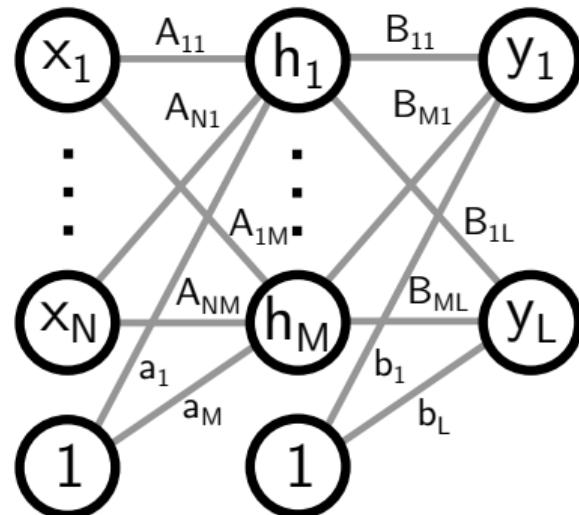
Linear Regression

► 3 Layers:

\mathbf{x} = Image

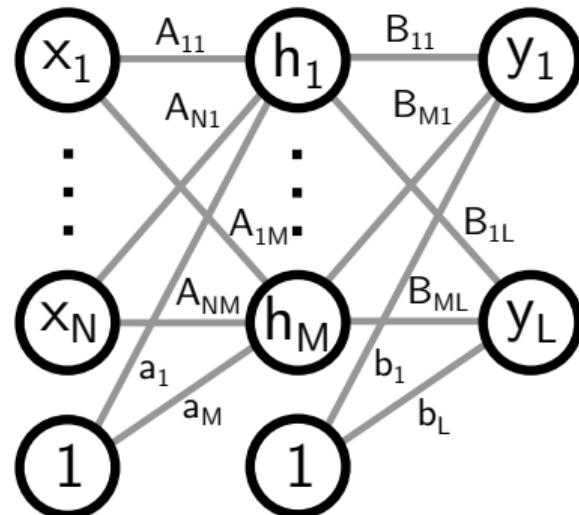
$\mathbf{h} = \mathbf{Ax} + \mathbf{a}$

$\mathbf{y} = \mathbf{Bh} + \mathbf{b}$



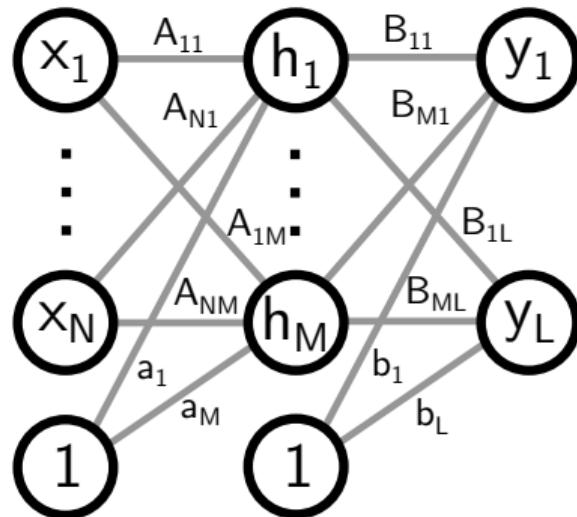
Linear Regression

- ▶ 3 Layers:
 - \mathbf{x} = Image
 - $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$
 - $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$
- ▶ Is this model better?



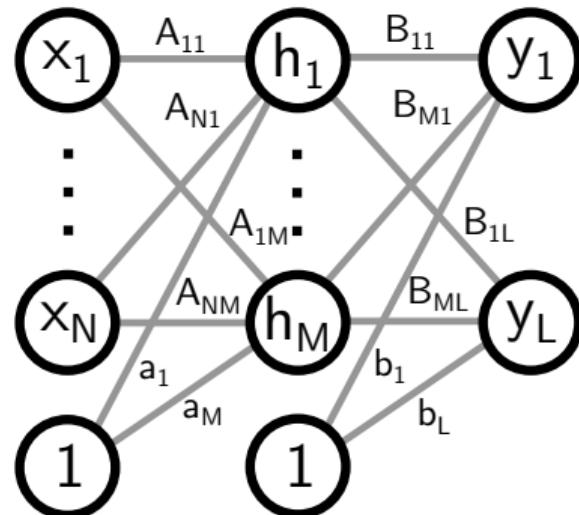
Linear Regression

- ▶ 3 Layers:
 \mathbf{x} = Image
 $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$
 $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$
- ▶ Is this model better?
 $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$



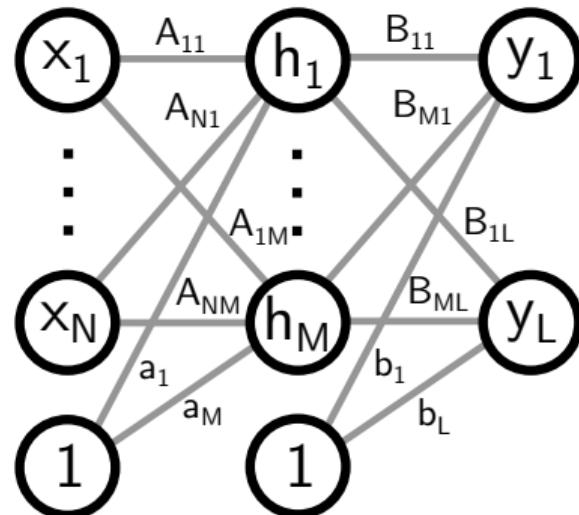
Linear Regression

- ▶ 3 Layers:
 - \mathbf{x} = Image
 - $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$
 - $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$
- ▶ Is this model better?
$$\mathbf{y} = \mathbf{B}(\mathbf{Ax} + \mathbf{a}) + \mathbf{b}$$



Linear Regression

- ▶ 3 Layers:
 \mathbf{x} = Image
 $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$
 $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$
- ▶ Is this model better?
 $\mathbf{y} = \mathbf{BAx} + \mathbf{Ba} + \mathbf{b}$



Linear Regression

- ▶ 3 Layers:

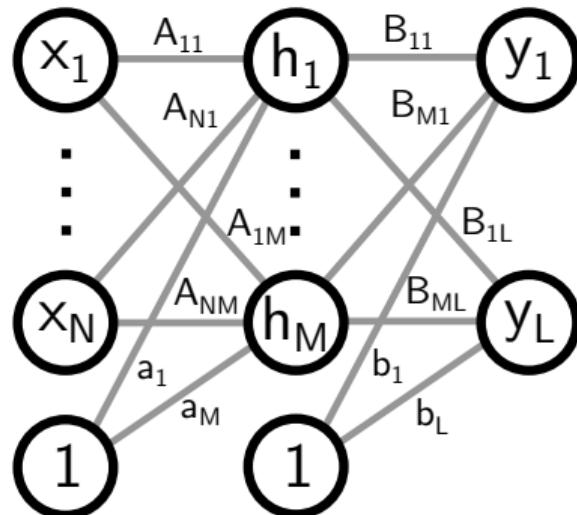
$$\mathbf{x} = \text{Image}$$

$$\mathbf{h} = \mathbf{Ax} + \mathbf{a}$$

$$\mathbf{y} = \mathbf{Bh} + \mathbf{b}$$

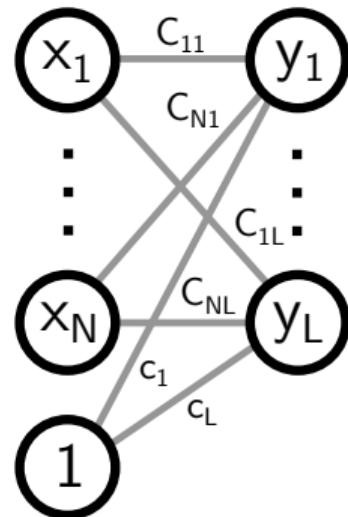
- ▶ Is this model better?

$$\mathbf{y} = \underbrace{\mathbf{BA}}_{=C} \mathbf{x} + \underbrace{\mathbf{Ba} + \mathbf{b}}_{=c}$$



Linear Regression

- ▶ 3 Layers:
 - \mathbf{x} = Image
 - $\mathbf{h} = \mathbf{Ax} + \mathbf{a}$
 - $\mathbf{y} = \mathbf{Bh} + \mathbf{b}$
- ▶ Is this model better?
 - $\mathbf{y} = \mathbf{Cx} + \mathbf{c}$



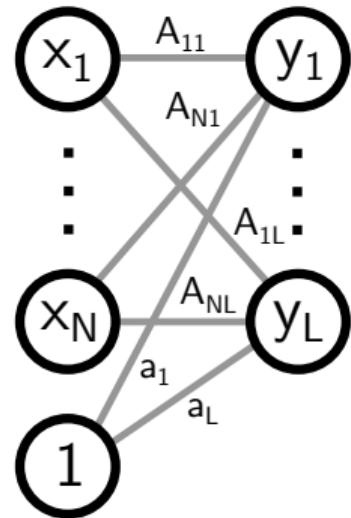
Logistic Regression

- ▶ Mapping:
 \mathbf{x} = Image



Logistic Regression

- ▶ Mapping:
 \mathbf{x} = Image
 $\mathbf{y} = \sigma(\mathbf{Ax} + \mathbf{a})$



Logistic Regression

- Mapping:

\mathbf{x} = Image

$\mathbf{y} = \sigma(\mathbf{Ax} + \mathbf{a})$

- With (elementwise):

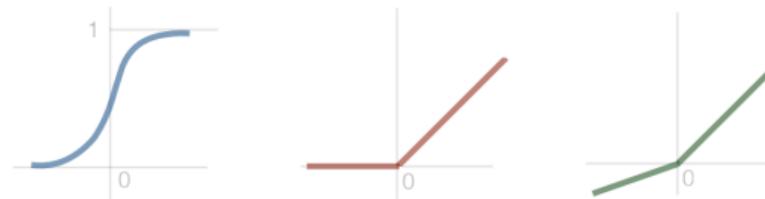
$$\sigma(x) = \frac{1}{1+\exp(-x)}$$

(a) Sigmoid

$$\sigma(z) = \frac{1}{1 + \exp(-z)} \quad \sigma(z) = \begin{cases} z, & z > 0 \\ 0, & z \leq 0 \end{cases} \quad \sigma(z) = \begin{cases} z, & z > 0 \\ az, & z \leq 0 \end{cases}$$

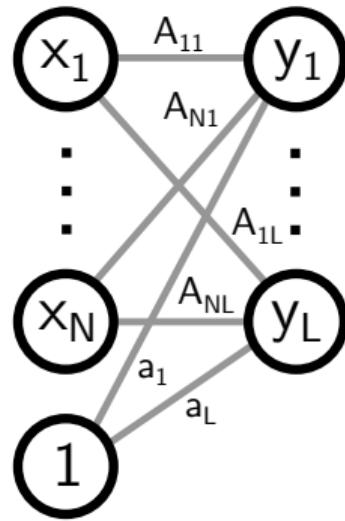
(b) ReLU

(c) PReLU



Logistic Regression

- ▶ Mapping:
 \mathbf{x} = Image
 $\mathbf{y} = \sigma(\mathbf{Ax} + \mathbf{a})$
- ▶ With (elementwise):
 $\sigma(x) = \frac{1}{1+\exp(-x)}$
- ▶ Classification:
 $L^* = \operatorname{argmax}_l y_l$



Multilayer Perceptron

- 3 Layers:

\mathbf{x} = Image

$$\mathbf{h} = \sigma(\mathbf{Ax} + \mathbf{a})$$

$$\mathbf{y} = \sigma(\mathbf{Bh} + \mathbf{b})$$



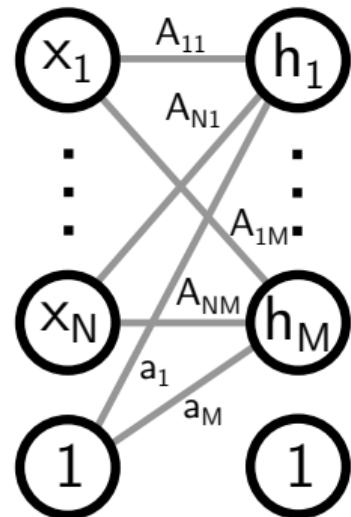
Multilayer Perceptron

- 3 Layers:

\mathbf{x} = Image

$$\mathbf{h} = \sigma(\mathbf{Ax} + \mathbf{a})$$

$$\mathbf{y} = \sigma(\mathbf{Bh} + \mathbf{b})$$



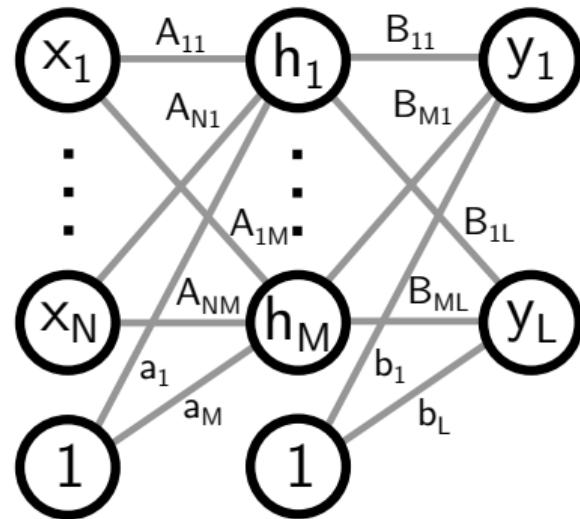
Multilayer Perceptron

- 3 Layers:

\mathbf{x} = Image

$$\mathbf{h} = \sigma(\mathbf{Ax} + \mathbf{a})$$

$$\mathbf{y} = \sigma(\mathbf{Bh} + \mathbf{b})$$



Multilayer Perceptron

- 3 Layers:

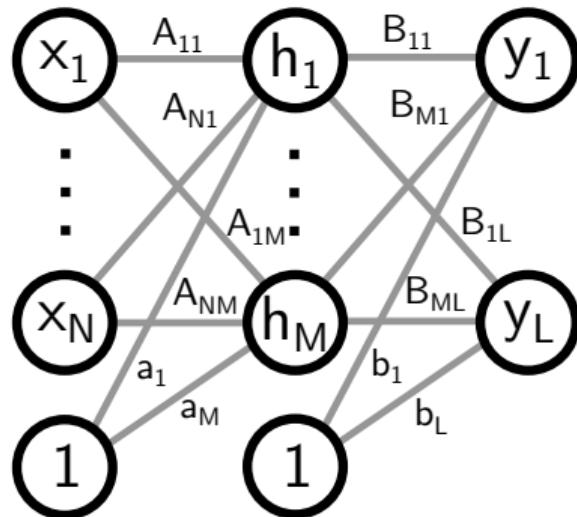
\mathbf{x} = Image

$$\mathbf{h} = \sigma(\mathbf{Ax} + \mathbf{a})$$

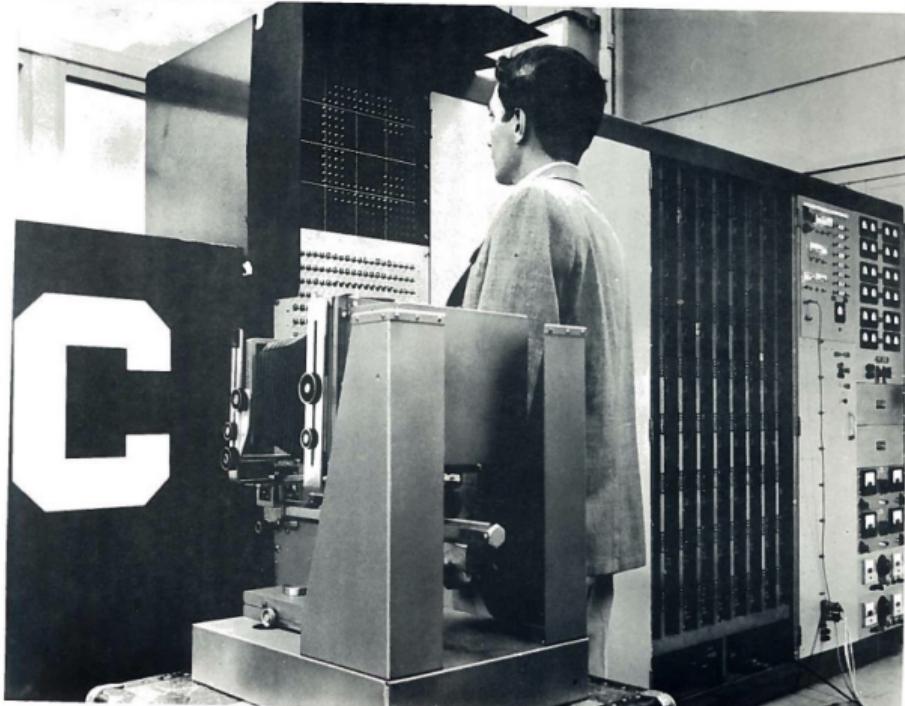
$$\mathbf{y} = \sigma(\mathbf{Bh} + \mathbf{b})$$

- Now:

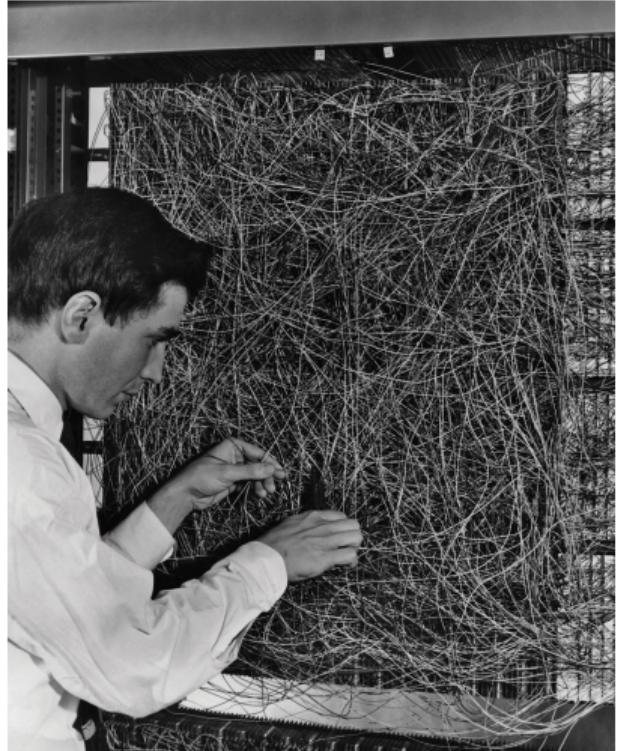
$$\mathbf{y} = \sigma(\mathbf{B}(\sigma(\mathbf{Ax} + \mathbf{a})) + \mathbf{b})$$



Perceptron: First Hardware Implementation (2 Layers)



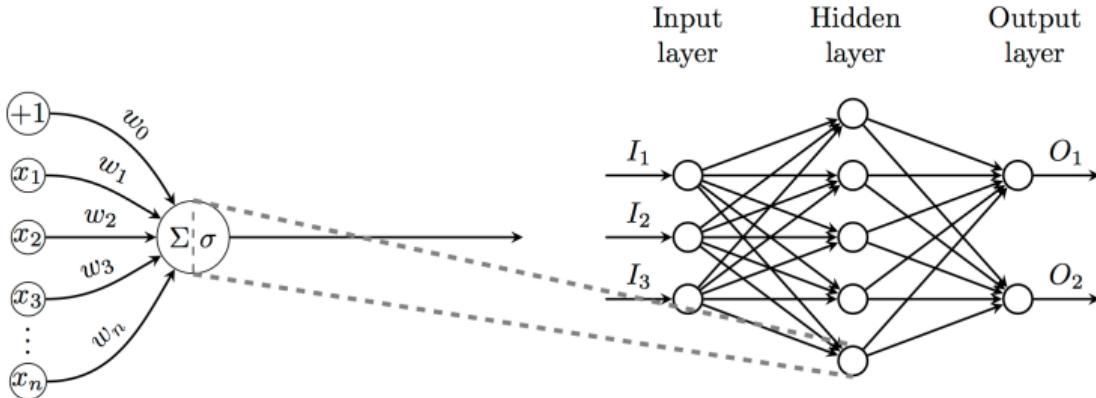
THE MARK I PERCEPTRON



Mark 1 Perceptron, Frank Rosenblatt, 1958

Rosenblatt: The Perceptron—a perceiving and recognizing automaton. Report 85-460-1, Cornell Aeronautical Laboratory, 1957.

Summary: Multilayer Perceptron in a Nutshell

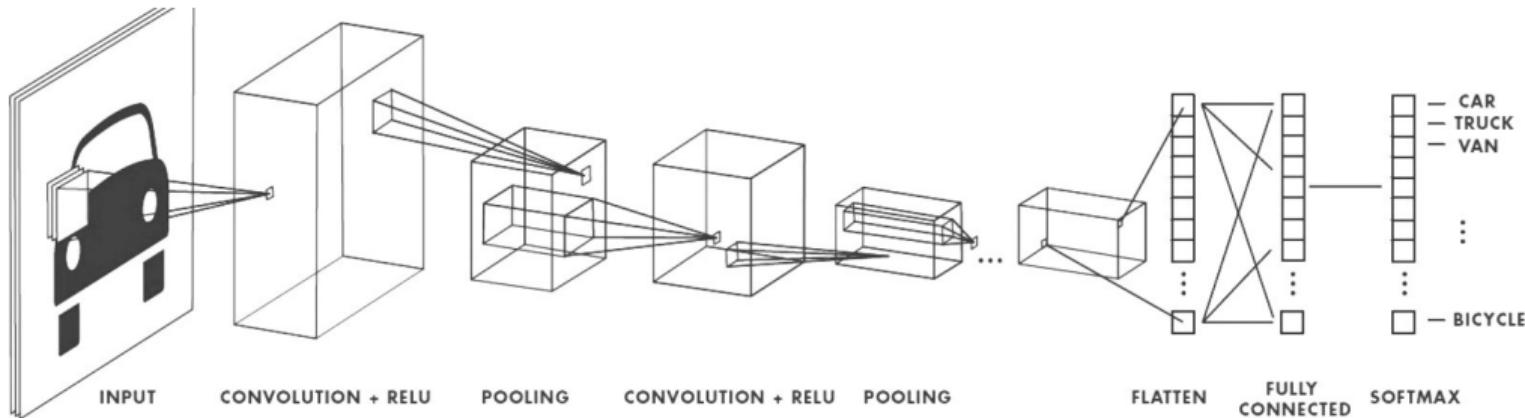


Layer j in a Multilayer Perceptron:

$$y_j = \sigma(x_j) = \underbrace{\frac{1}{1 + \exp\{-x_j\}}}_{\text{activation function}}$$
$$x_j = \underbrace{\sum_i w_{ji} y_i}_{\text{weighted sum}}$$

- ▶ Linearly combines *all* outputs y_i of previous layer i into *one* weighted sum x_j
 - ▶ Thus the layers in a MLP are also called "fully connected layers"
- ▶ Applies non-linear activation function $\sigma(x_j)$, yielding new output y_j

Convolutional Neural Networks



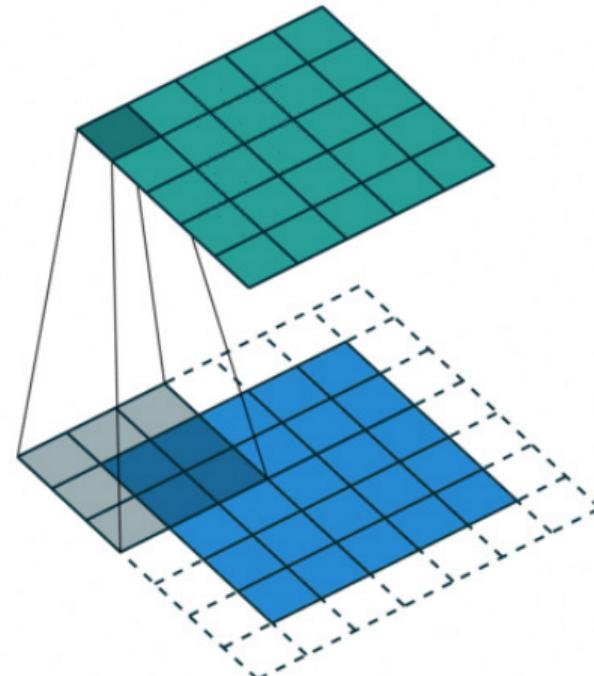
- ▶ Images are typically very high-dimensional. What is the problem with MLPs?
- ▶ ConvNets represent data in 3 dimensions: width, height, depth (= feature maps)
- ▶ ConvNets interleave discrete convolutions, non-linearities and pooling
- ▶ The learned weights now correspond to the convolution kernel parameters
- ▶ Key ideas: sparse interactions, parameter sharing, equivariant representation
- ▶ Biological example: simple cells and complex cells in primary visual cortex (V1)

Building Block: Convolution Layer

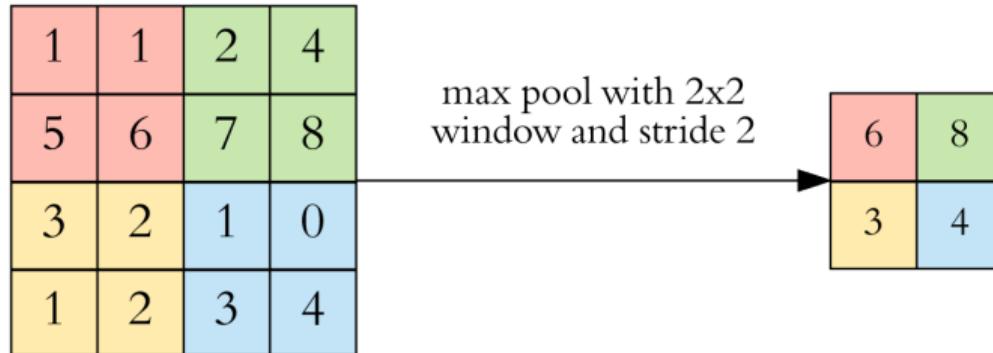
Convolution Operation:

$$x_j(u, v, \cdot) = \sum_{u', v', d} k(u', v', d) y_i(u + u', v + v', d)$$

- ▶ With pixel (u, v) , channel d , kernel $k(u, v, d)$
- ▶ Should be called “cross-correlation”
- ▶ Followed by non-linearity (e.g., ReLU)
- ▶ Right: Example for convolution with stride 1
 - ▶ Here: one feature map as input but in practice multiple feature maps
- ▶ Let's do a little quiz!



Building Block: Pooling Layer



- ▶ Combine multiple inputs locally into one number
- ▶ Most common pooling operations: maximum, minimum, average
- ▶ Stride 2 increases receptive field size by factor 2 in each dimension
- ▶ How many trainable parameters/weights?
- ▶ Convolutions can also be strided!

LeNet-5

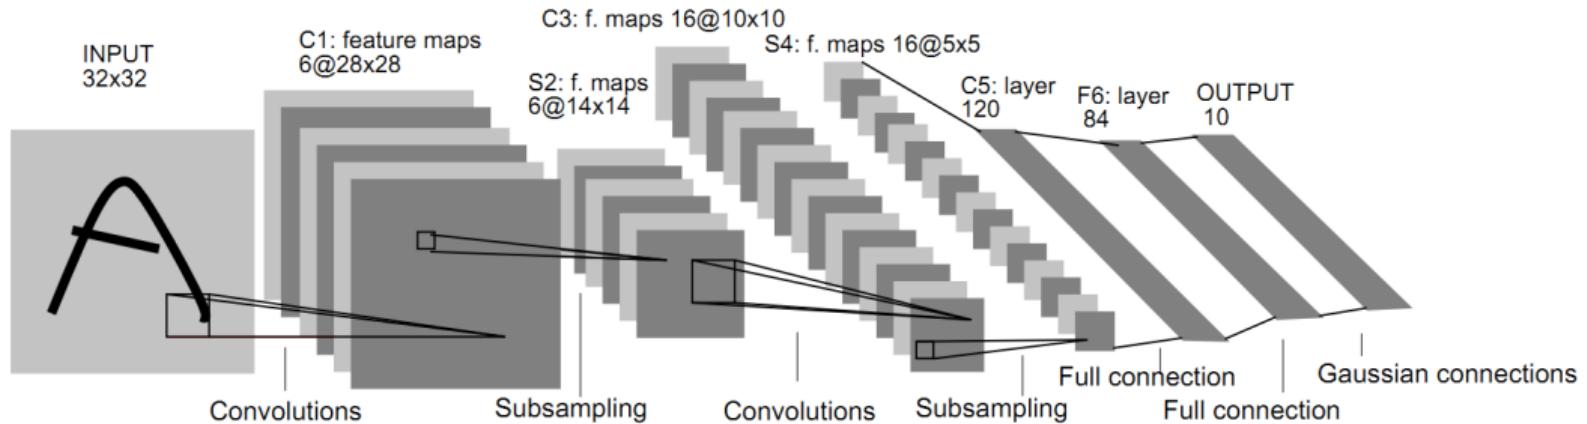


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

- ▶ First implementation of convolutional net with learned weights
- ▶ 2 convolution layers (5x5), 2 subsampling layers (stride 2), 2 fully connected layers

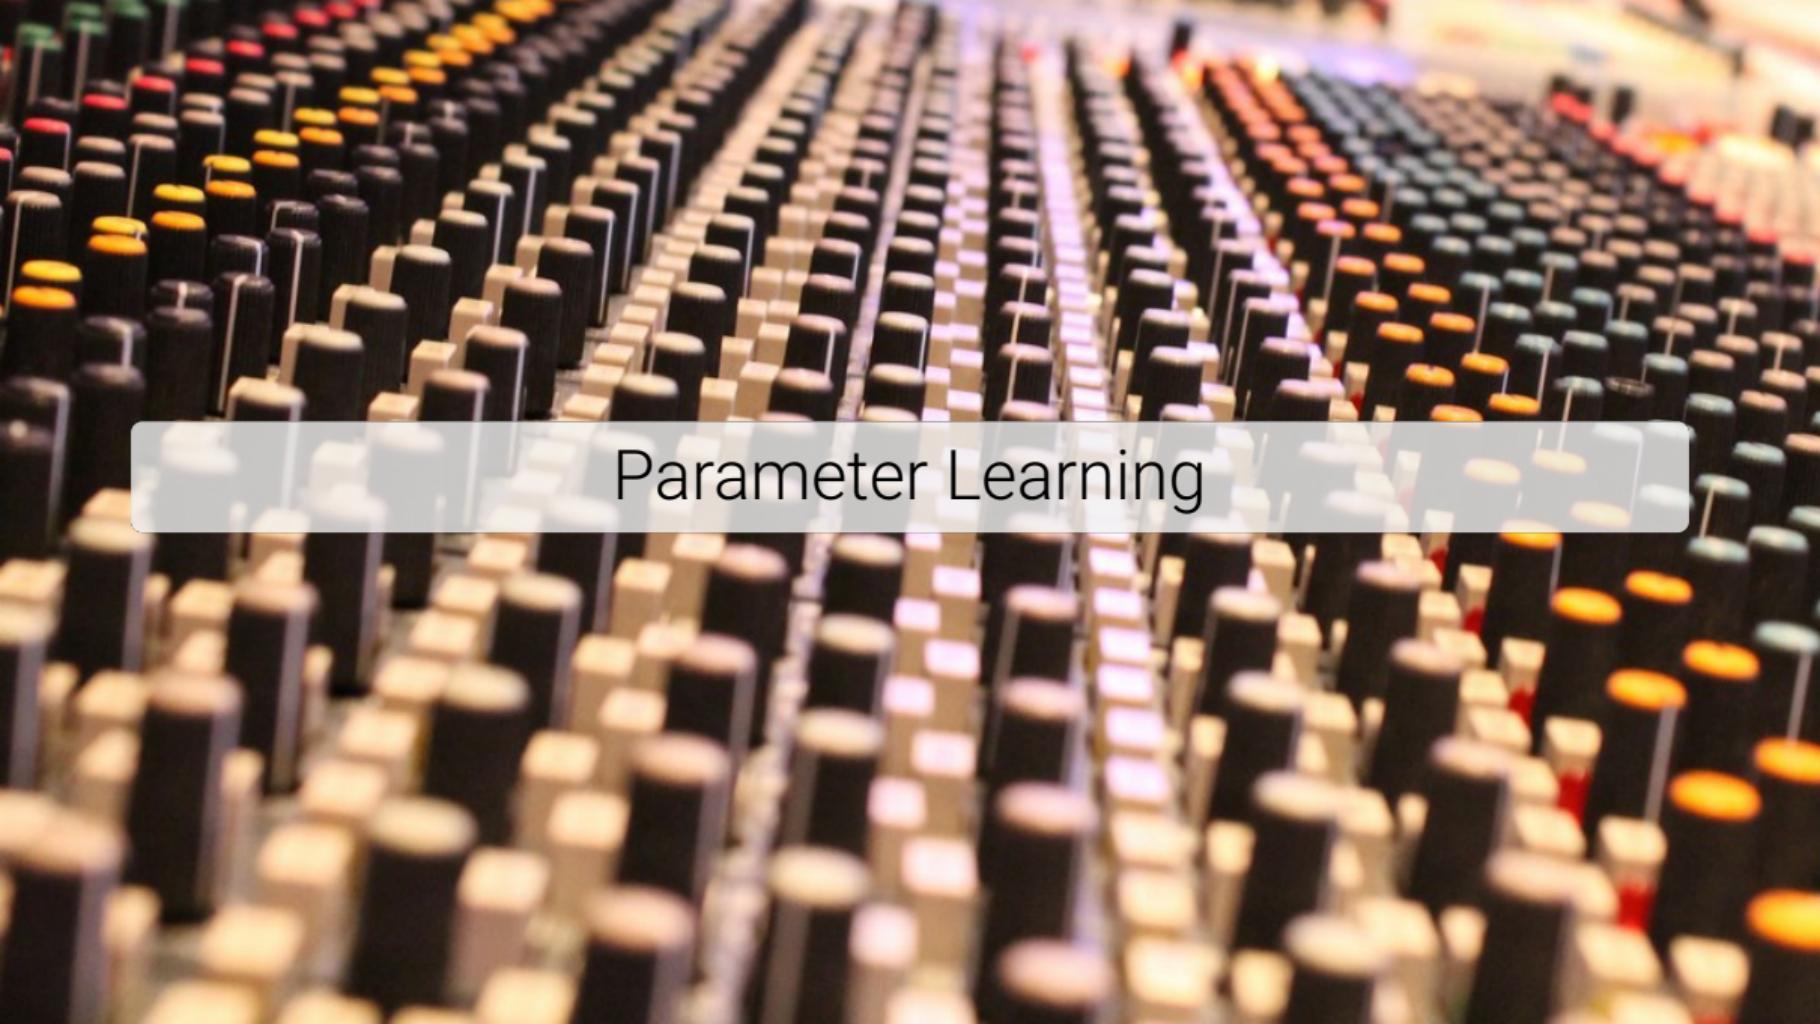
MNIST Digit Classification



VGG-19

- ▶ Evaluating networks of increasing depth (today: 100-1000 layers)
- ▶ Demonstrated significant improvement by increasing the depth to 16-19 weight layers
- ▶ Key to keep number of parameters tractable: Small 3×3 filters in all convolutional layers (stride 1)

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					



Parameter Learning

Multilayer Perceptron - Learning

Learning:

- ▶ Define error function over all data points $\{(\mathbf{x}_n, \mathbf{t}_n)\}_{n=1}^N$, e.g.:

$$E = \frac{1}{2} \sum_n \sum_j (y_{jn} - t_{jn})^2$$

- ▶ Here y_{jn} is an element of the last layer j of the network
- ▶ Calculate $\nabla_{\mathbf{W}} E = \sum_n \nabla_{\mathbf{W}} E_n$ with $E_n = \frac{1}{2} \sum_j (y_{jn} - t_{jn})^2$
- ▶ Gradient descent: $\mathbf{W}^{t+1} = \mathbf{W}^t - \eta \nabla_{\mathbf{W}} E|_{\mathbf{W}^t}$ (η = "learning rate")

Problem?

- ▶ Gradient needs to be computed for all data points (problematic when N is large)
 - ▶ Solution: Perform Stochastic Gradient Descent (SGD) using mini batches:
 - ▶ Approximate gradient with $M \ll N$ randomly picked data points at each iteration (typically $M = 16, \dots, 128$ as gradients too noisy for $M = 1$)
- ▶ In other words: $\nabla_{\mathbf{W}} E \approx \sum_m \nabla_{\mathbf{W}} E_m$. But how to calculate $\nabla_{\mathbf{W}} E_m$?

Backpropagation

[Rumelhart et al., Nature 1986]

Backpropagation on two Slides

Neural Network Layer:

$$y_j = \frac{1}{1 + \exp\{-x_j\}} \quad x_j = \sum_i w_{ji} y_i$$

Goal: Calculate gradient $\nabla_{\mathbf{W}} E_m = \nabla_{\mathbf{W}} \frac{1}{2} \sum_j (y_{jm} - t_{jm})^2$ (dropping m for simplicity)

$$\frac{\partial E}{\partial y_j} = y_j - t_j$$

$$\underbrace{\frac{\partial E}{\partial x_j} = \frac{\partial E}{\partial y_j} \frac{\partial y_j}{\partial x_j}}_{=: \delta_j} = (y_j - t_j) y_j (1 - y_j)$$

$$\frac{\partial E}{\partial w_{ji}} = \frac{\partial E}{\partial x_j} \frac{\partial x_j}{\partial w_{ji}} = \delta_j y_i$$

$$\frac{\partial E}{\partial y_i} = \sum_j \frac{\partial E}{\partial x_j} \frac{\partial x_j}{\partial y_i} = \sum_j \delta_j w_{ji}$$

$$\frac{\partial E}{\partial x_i} = \frac{\partial E}{\partial y_i} \frac{\partial y_i}{\partial x_i} = \boxed{\delta_i := y_i(1 - y_i) \sum_j \delta_j w_{ji}}$$

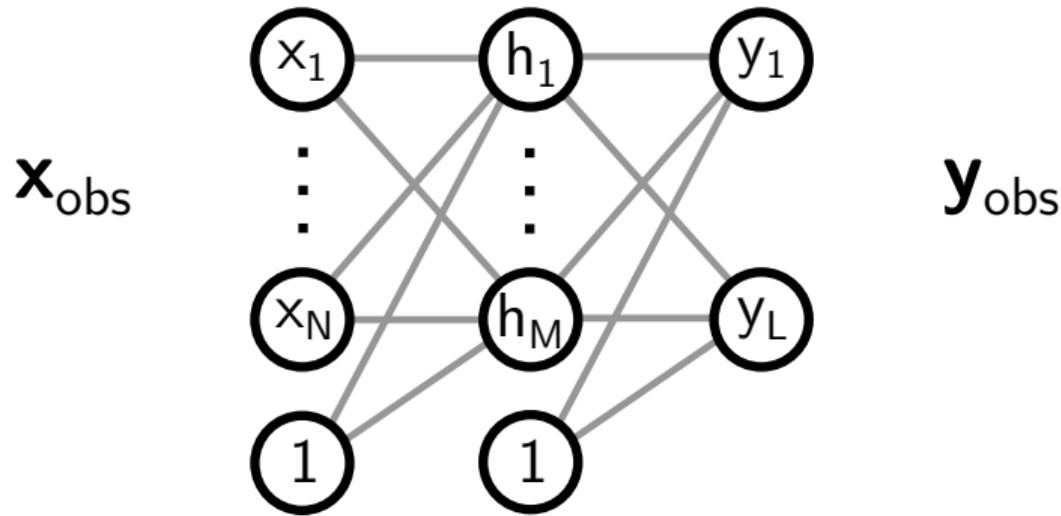
Backpropagation formula

Backpropagation on two Slides

Algorithm for training a neural network using SGD:

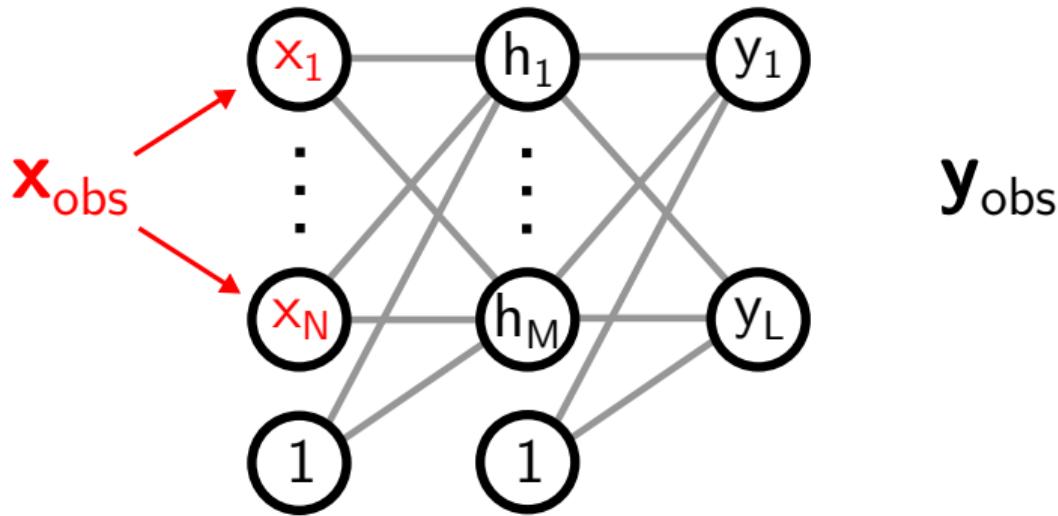
1. Randomly initialize \mathbf{W}
2. Sample mini-batch $\mathcal{X}_{\text{batch}} \subset \mathcal{X}$
3. For each data point $\mathbf{x}_m \in \mathcal{X}_{\text{batch}}$ do:
 - 3.1 **Forward propagate** \mathbf{x}_m through network: $y_j = \sigma(\sum_i w_{ji}y_i)$ (here: $\{y_0\} = \mathbf{x}_m$)
 - 3.2 **Calculate errors** δ_j at output nodes: $\delta_j = (y_j - t_j)y_j(1 - y_j)$
 - 3.3 **Backpropagate errors** δ_i through network: $\delta_i = y_i(1 - y_i) \sum_j \delta_j w_{ji}$
 - 3.4 **Calculate derivatives** $\frac{\partial E_m}{\partial w_{ji}} = \delta_j y_i$
4. Update gradients $\mathbf{W}^{t+1} = \mathbf{W}^t - \eta \sum_{m=1}^M \nabla_{\mathbf{W}} E_m |_{\mathbf{W}^t}$
5. If validation error decreases return to step 2, otherwise stop

Illustration: Forward Pass



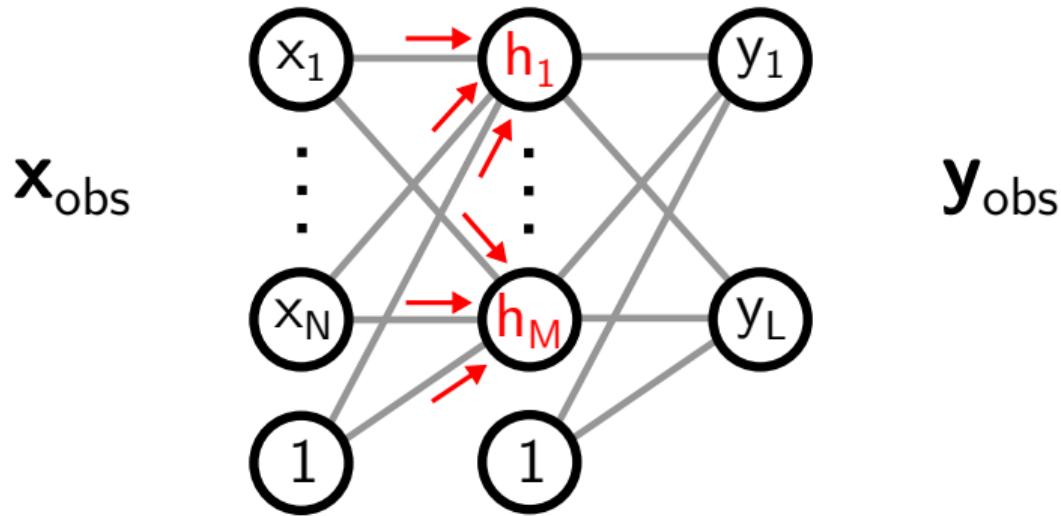
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Forward Pass



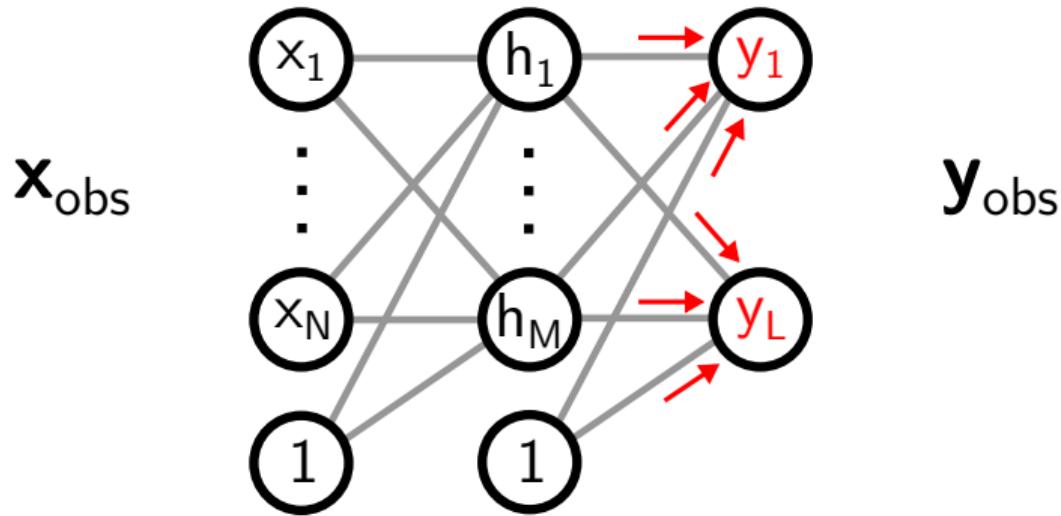
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Forward Pass



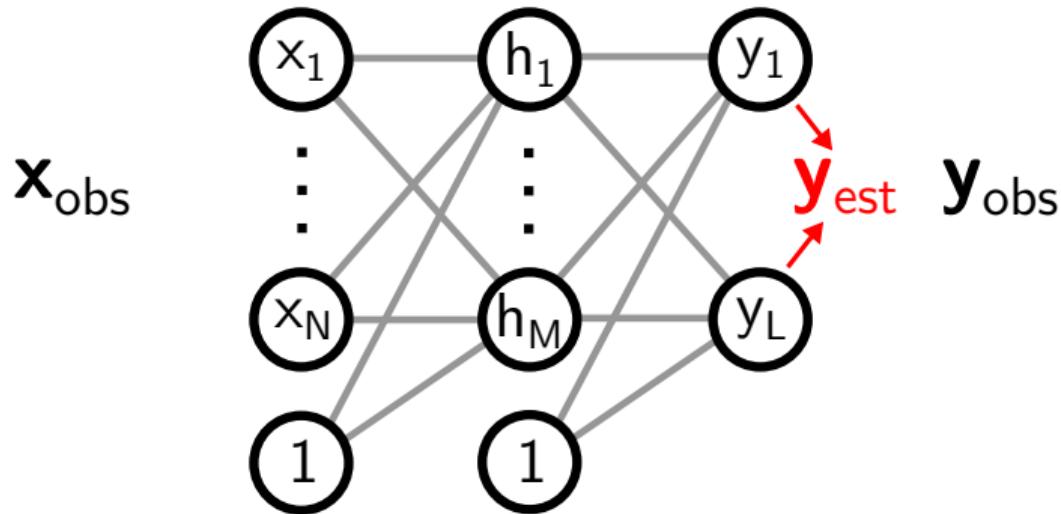
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Forward Pass



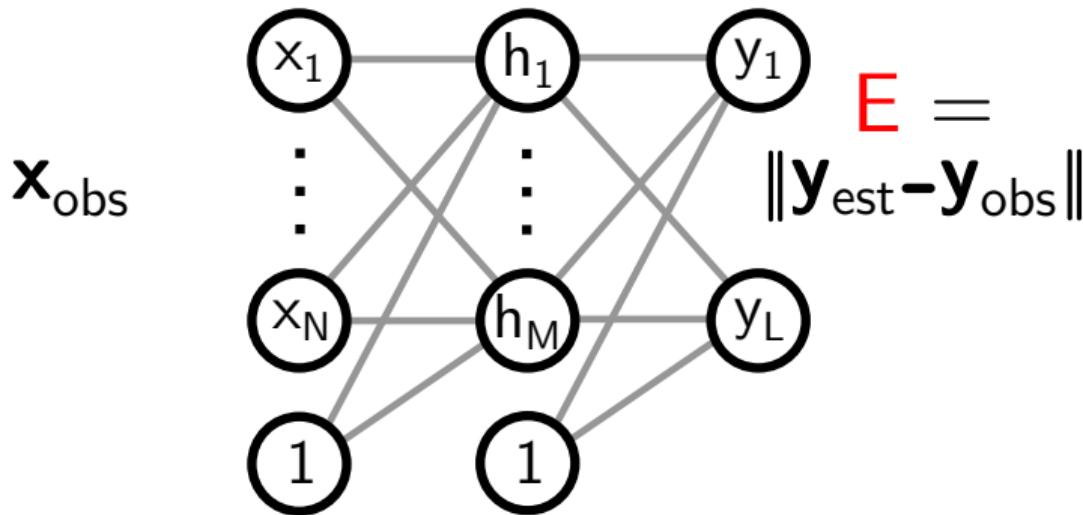
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Forward Pass



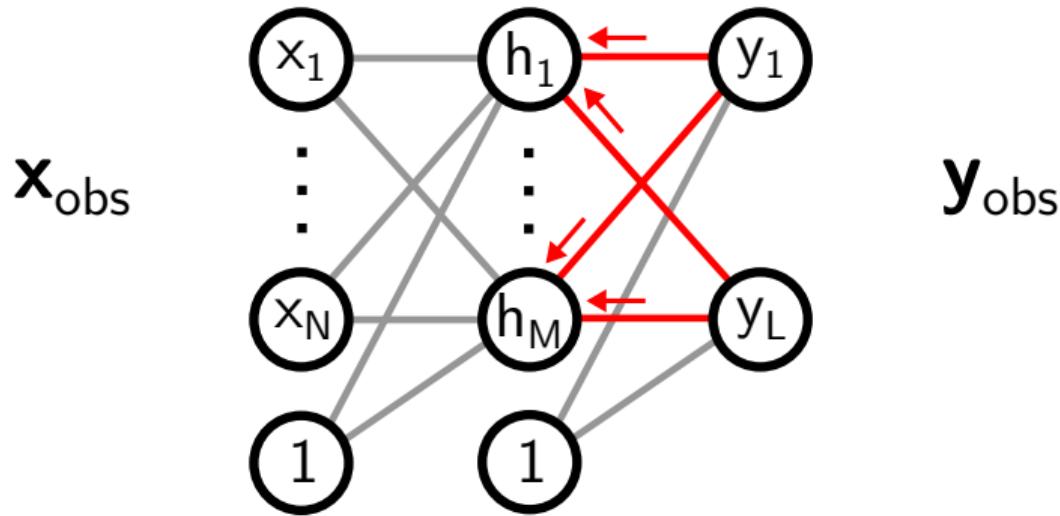
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Error Calculation



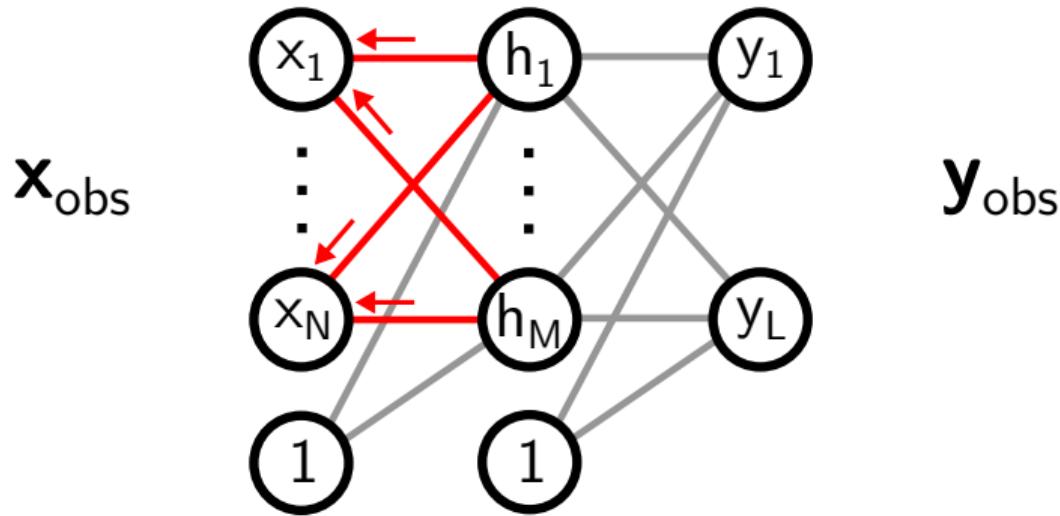
- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Error Backpropagation



- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Illustration: Error Backpropagation



- ▶ Forward pass: resembles processing in the brain
- ▶ Backward pass: relationship to neural learning less clear / less well understood

Comments on Backpropagation

Deep neural networks are usually trained using backpropagation.

- ▶ Unfortunately, neural nets are non-convex, many local minima
- ▶ Can get stuck in poor local minima, but: more parameters can help!
- ▶ Careful design needed
- ▶ Monitor loss, training and validation error

Overfitting remains a problem due to the large number of parameters.

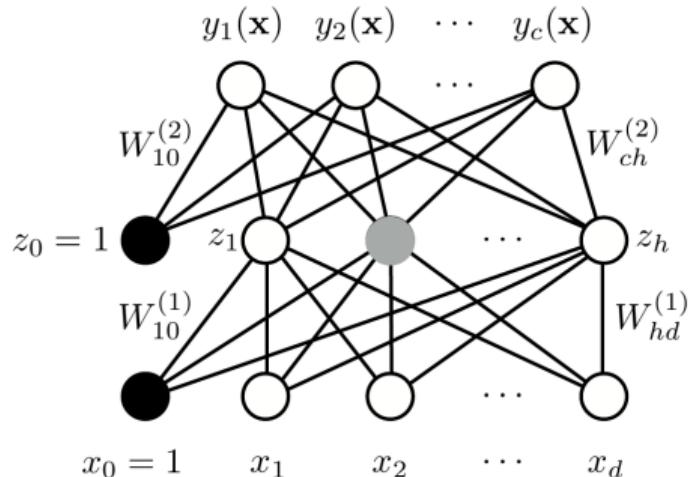
Measures to reduce overfitting?

- ▶ Weight decay (regularize weights): $E = E + \lambda \sum_j \sum_i w_{ji}^2$
- ▶ Drop-out (next slide)
- ▶ Data augmentation (next slide)

Drop-out

Idea:

- ▶ During forward pass: set every hidden unit to zero with probability 0.5
- ▶ Backward pass as before
- ▶ Avoids co-adaptation, introduces randomness, simulates network ensemble

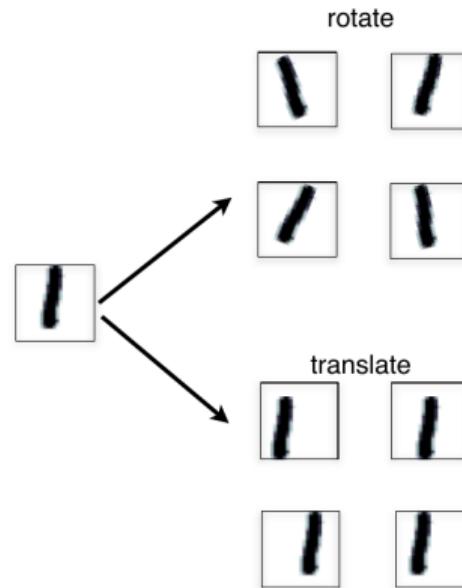


Data Augmentation

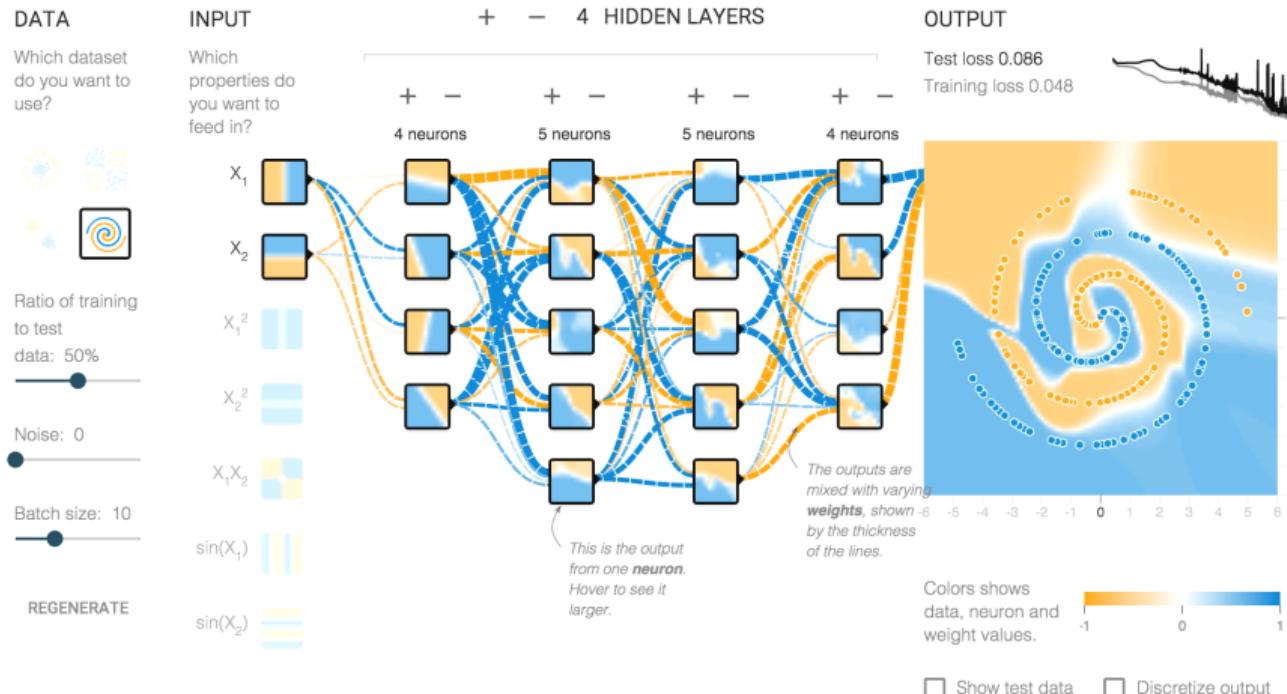
Idea:

- ▶ Augment training set by “jittering”
- ▶ Transformation must be label preserving

1 1 5 4 3
7 5 3 5 3
5 5 9 0 6
3 5 2 0 0



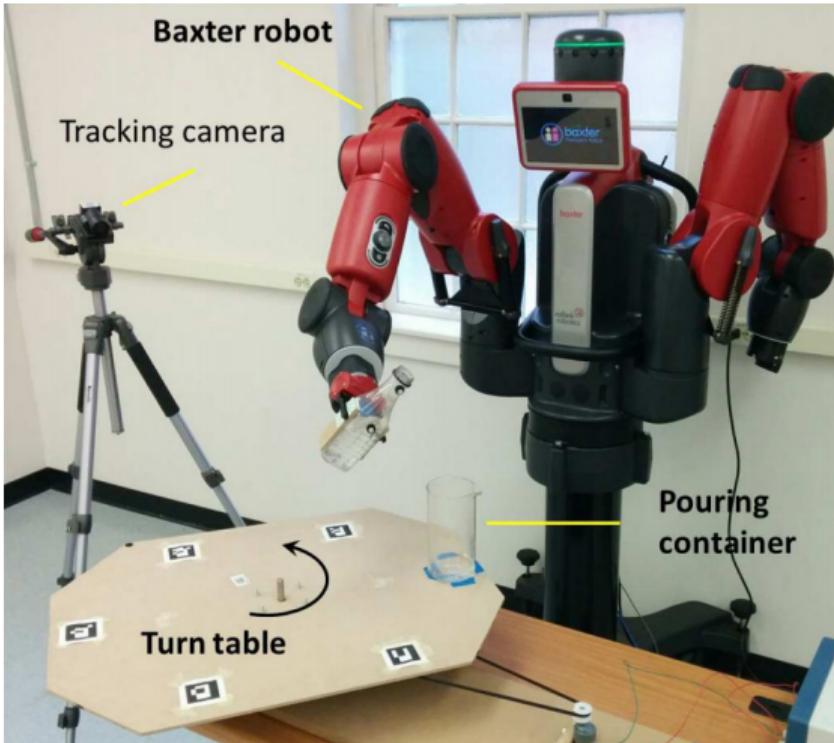
Tensorflow Playground



<http://playground.tensorflow.org>

Imitation Learning

Imitation Learning: Manipulation



Imitation Learning: Ghosting

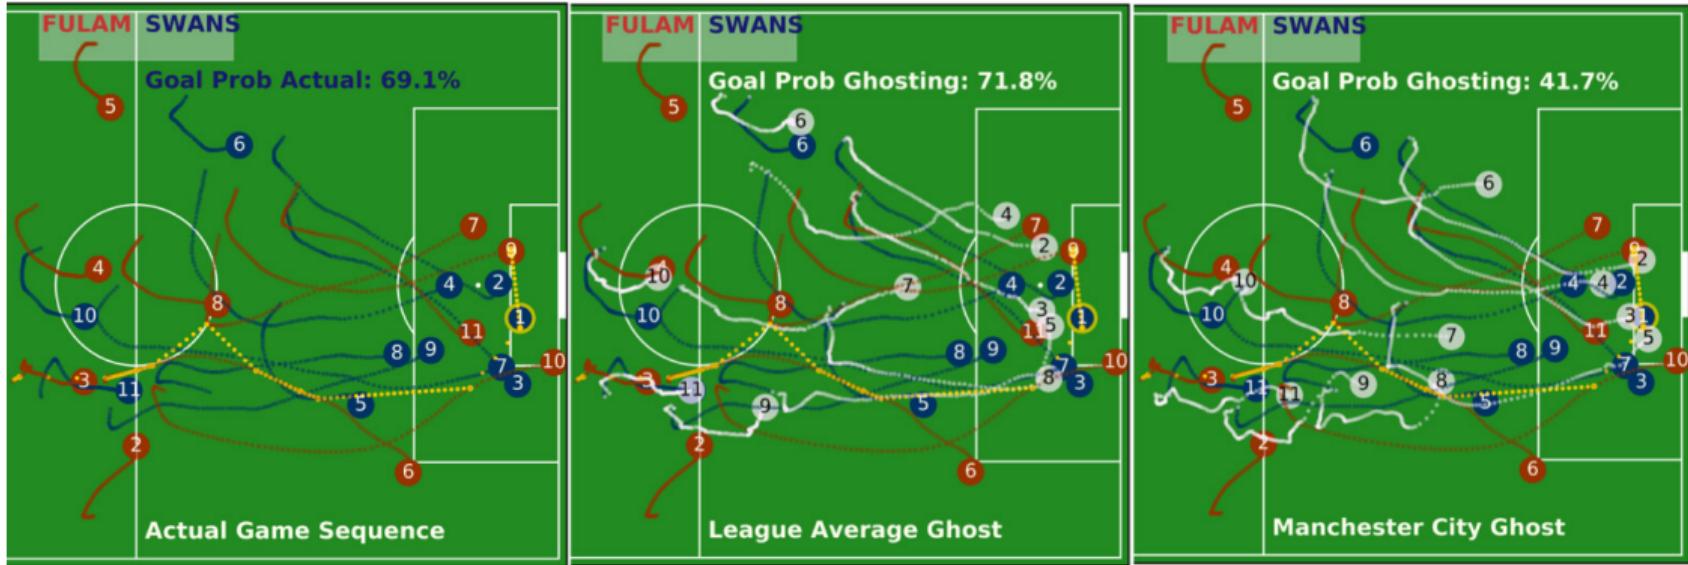
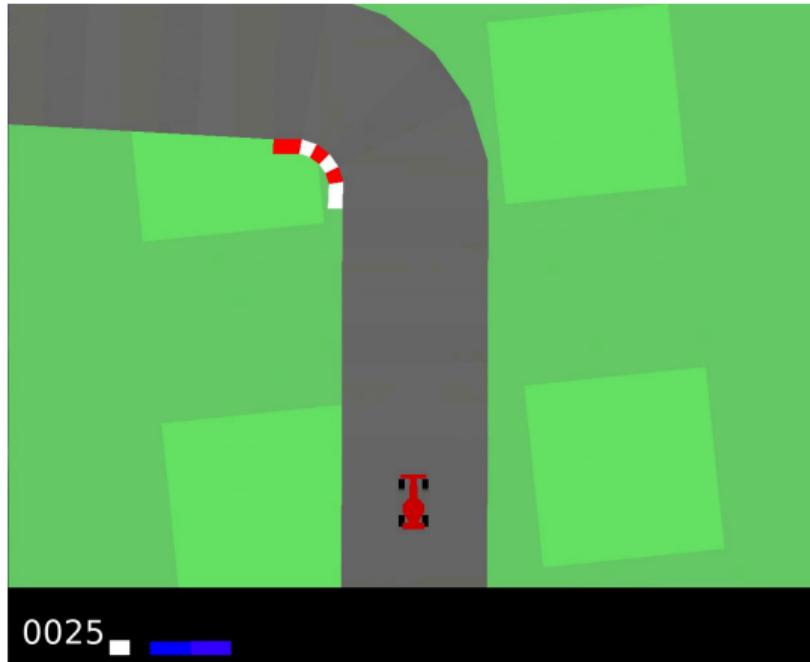


Figure 1 Our data-driven ghosting method can be applied to various game contexts to better understand defensive strategies. In the depicted scenario, Fulham (red) scores a goal on Swansea (blue). Ghosts (white) represent where Swansea defenders should have been according to a league average model (LAG) and Manchester City model (MUG).

Imitation Learning: Car Racing

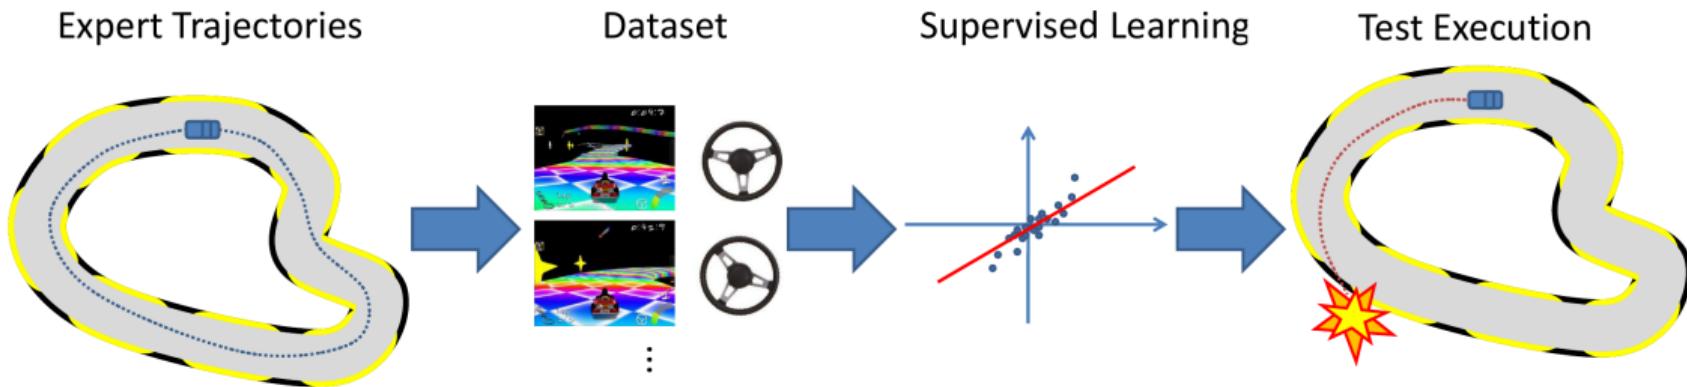


Trainer
(Human Driver)



Trainee
(Neural Network)

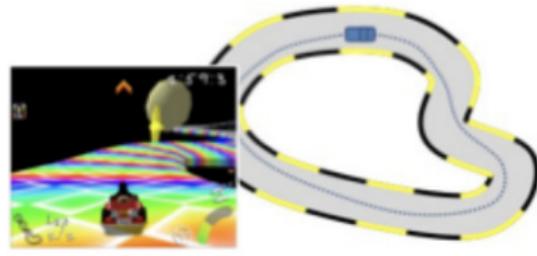
Imitation Learning in a Nutshell



Hard coding policies is often difficult \Rightarrow Rather use a data-driven approach!

- ▶ **Given:** demonstrations or demonstrator
- ▶ **Goal:** train a policy to mimic decision
- ▶ **Variants:** behavior cloning (this lecture), inverse optimal control, ...

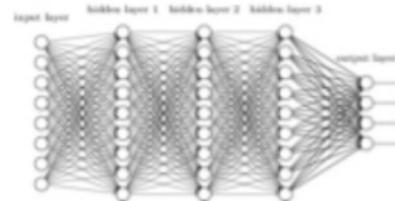
Ingredients of Imitation Learning



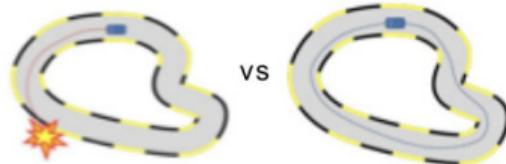
Demonstrations or Demonstrator



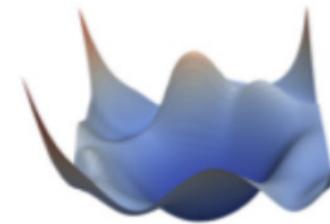
Environment / Simulator



Policy Class



Loss Function



Learning Algorithm

Formal Definition of Imitation Learning

Components:

- ▶ State: $s \in \mathcal{S}$ may be partially observed (e.g., game screen)
- ▶ Action: $a \in \mathcal{A}$ may be discrete or continuous (e.g., turn angle, speed)
- ▶ Policy: $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ we want to learn the policy parameters θ
- ▶ Optimal action: $a^* \in \mathcal{A}$ provided by expert demonstrator
- ▶ Optimal policy: $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ provided by expert demonstrator
- ▶ State dynamics: $P(s_{i+1}|s_i, a_i)$ simulator, typically not known to policy
Often deterministic: $s_{i+1} = T(s_i, a_i)$ deterministic mapping
- ▶ Rollout: Given s_0 , sequentially execute $a_i = \pi_\theta(s_i)$ & sample $s_{i+1} \sim P(s_{i+1}|s_i, a_i)$
yields trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$ let's do an example!
- ▶ Loss function: $\mathcal{L}(a^*, a)$ loss of action a given optimal action a^*

Formal Definition of Imitation Learning

General Imitation Learning:

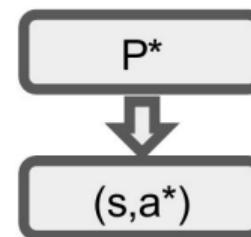
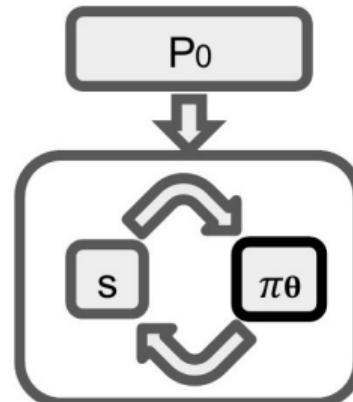
$$\operatorname{argmin}_{\theta} \mathbb{E}_{s \sim P(s|\pi_{\theta})} [\mathcal{L}(\pi^*(s), \pi_{\theta}(s))]$$

- ▶ State distribution $P(s|\pi_{\theta})$ depends on rollout determined by current policy π_{θ}

Behavior Cloning:

$$\operatorname{argmin}_{\theta} \underbrace{\mathbb{E}_{(s^*, a^*) \sim P^*} [\mathcal{L}(a^*, \pi_{\theta}(s^*))]}_{= \sum_{i=1}^N \mathcal{L}(a_i^*, \pi_{\theta}(s_i^*))}$$

- ▶ State distribution P^* provided by expert
- ▶ Reduces to supervised learning problem

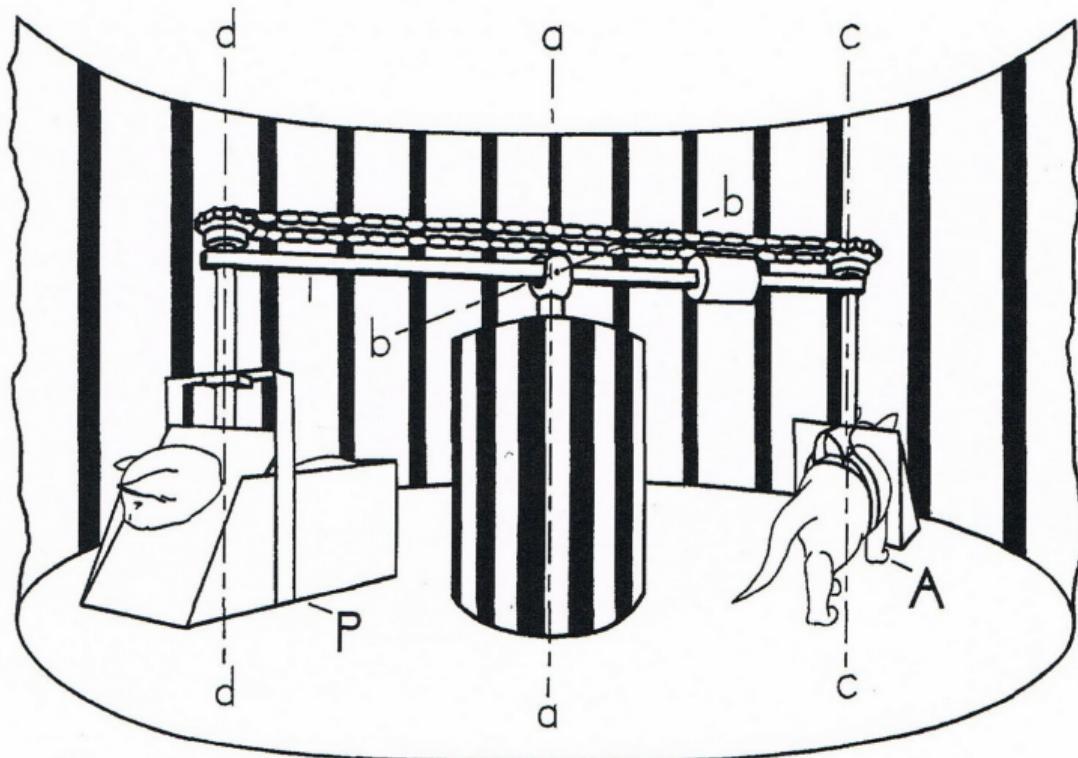


What is the problem?

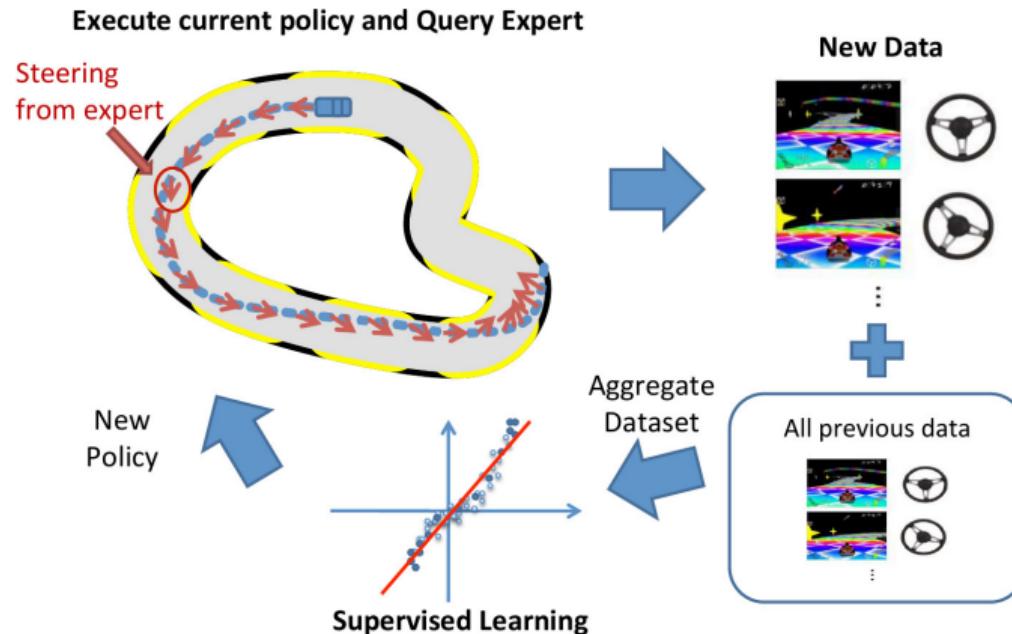
Challenges of Behavior Cloning

- ▶ Behavior cloning reasons only about immediate next step
- ▶ Behavior cloning makes IID assumption
 - ▶ Next state is sampled from states observed during expert demonstration
 - ▶ Thus, next state is sampled independently from action predicted by current policy
- ▶ What if π_θ makes a mistake?
 - ▶ Enters new states that haven't been observed before
 - ▶ New states not sampled from same (expert) distribution anymore
 - ▶ Cannot recover, catastrophic failure in the worst case
- ▶ What can we do to overcome this train/test distribution mismatch?

Experiment by Held and Hein



DAGGER: Dataset Aggregation



- **Idea:** Iteratively build a set of inputs that the final policy is likely to encounter based on previous experience. Query expert for aggregate dataset.

DAGGER: Dataset Aggregation

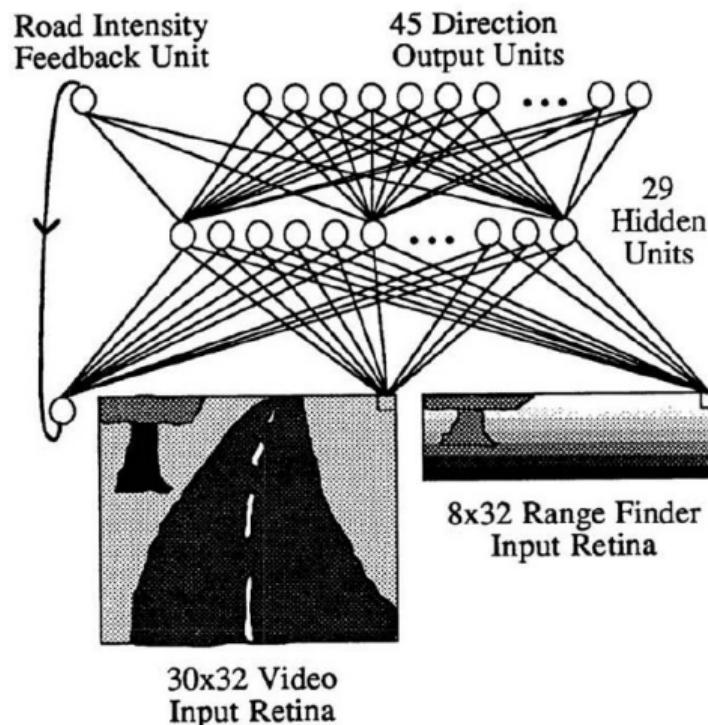
```
Initialize  $\mathcal{D} \leftarrow \emptyset$ .
Initialize  $\hat{\pi}_1$  to any policy in  $\Pi$ .
for  $i = 1$  to  $N$  do
    Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .
    Sample  $T$ -step trajectories using  $\pi_i$ .
    Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$ 
    and actions given by expert.
    Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .
    Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .
end for
Return best  $\hat{\pi}_i$  on validation.
```

- ▶ Let T be the number of rollout time steps and ϵ be the probability of a mistake
- ▶ Behavior cloning: $O(T^2\epsilon)$ regret, DAGGER: no regret (in suitable conditions)
 - ▶ Regret = loss of current policy - loss of perfect policy
- ▶ But requires to flexibly run the environment and query the expert during training

ALVINN: An Autonomous Land Vehicle in a Neural Network

ALVINN: An Autonomous Land Vehicle in a Neural Network

- ▶ Fully connected 3 layer neural net
- ▶ 36k parameters
- ▶ Maps road images to turn radius
- ▶ Directions discretized (45 bins)
- ▶ Trained on simulated road images!
- ▶ Tested on unlined paths, lined city streets and interstate highways
- ▶ 90 consecutive miles at up to 70 mph

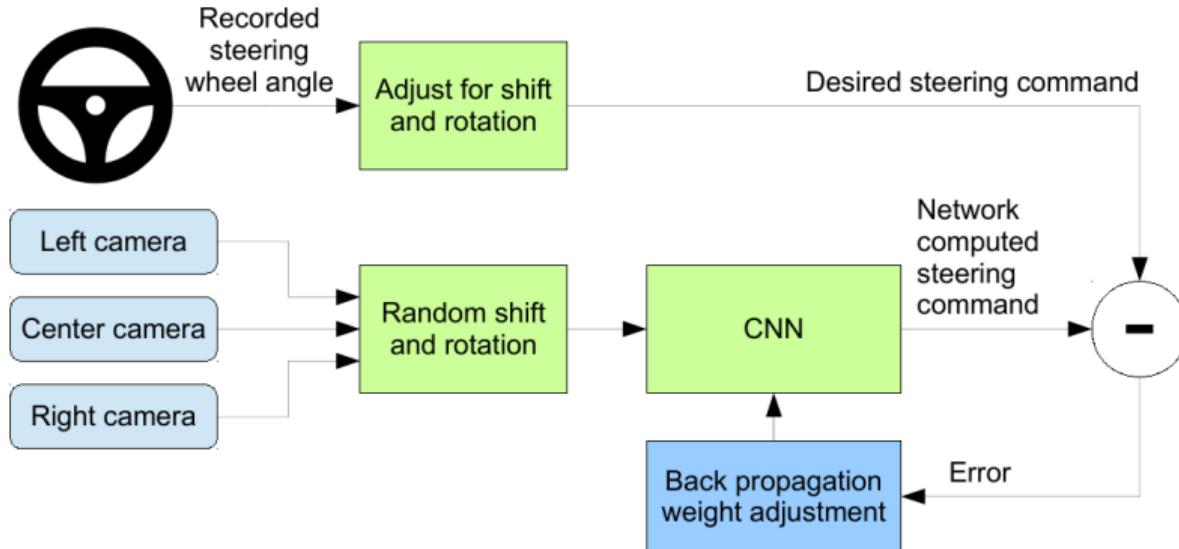


ALVINN: An Autonomous Land Vehicle in a Neural Network



End-to-End Learning for Self-Driving Cars

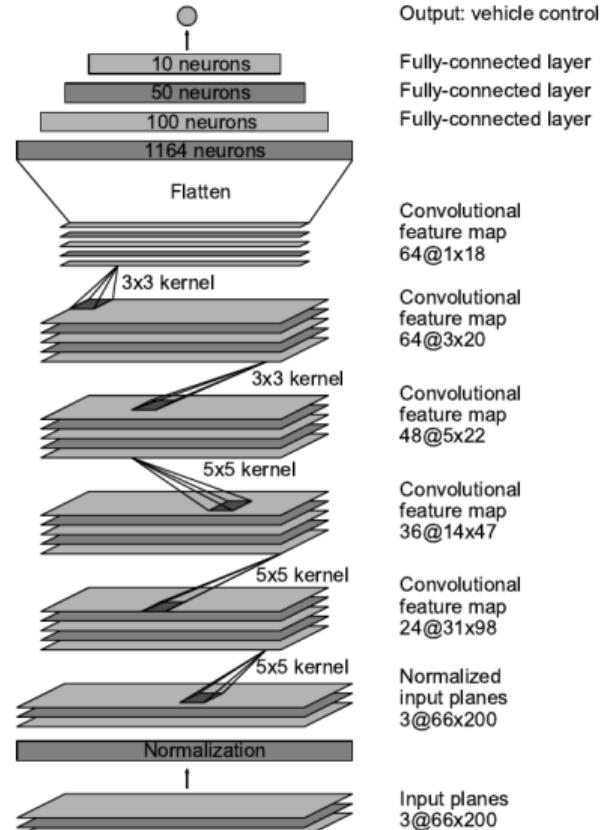
PilotNet: System Overview



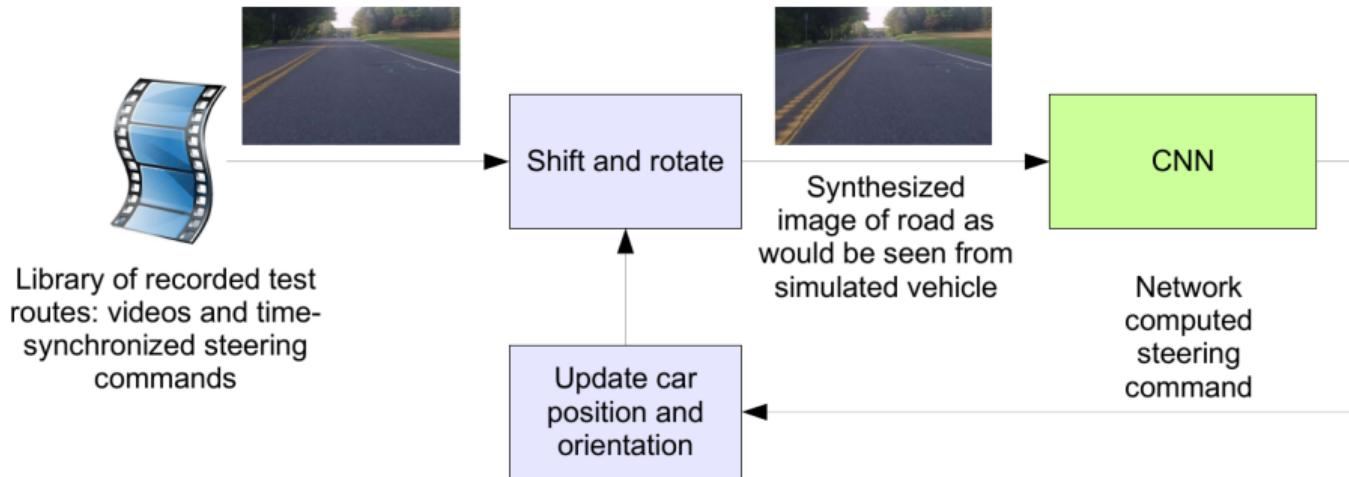
- ▶ Steering command represented as $1/r$ where r = turning radius in meters
- ▶ Data augmentation by 3 cameras and virtually shifted / rotated images assuming the world is flat (homography), adjusting the steering angle appropriately

PilotNet: Architecture

- ▶ Convolutional network (250k param)
- ▶ Input: YUV image representation
- ▶ 1 Normalization layer
 - ▶ Not learned
- ▶ 5 Convolutional Layers
 - ▶ 3 strided 5x5
 - ▶ 2 non-strided 3x3
- ▶ 3 Fully connected Layers
- ▶ Output: turning radius
- ▶ No lateral control!
- ▶ Trained on 72h of driving



PilotNet: Simulation



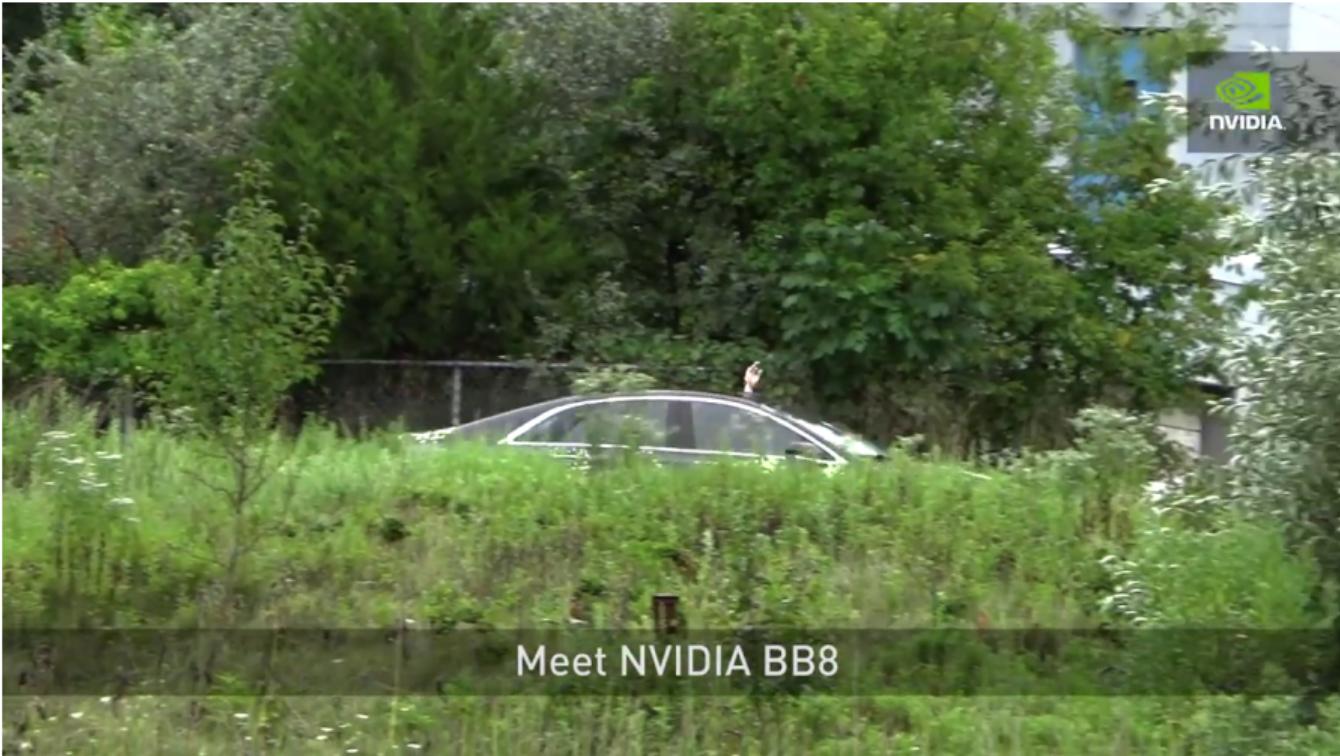
- ▶ Simulator takes pre-recorded videos and extracts lane center
- ▶ Creates novel viewpoints and applies CNN commands to vehicle model
- ▶ Updates the next image based on the resulting vehicle pose

PilotNet: Simulation



- ▶ Interventions (1m deviation) subtract 6 seconds from autonomy measure

PilotNet: Video



Further Readings

- ▶ Ross, Gordon and Bagnell: A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning. AISTATS, 2011.
- ▶ Pomerleau: ALVINN: An Autonomous Land Vehicle in a Neural Network. NIPS, 1988.
- ▶ Bojarski et al.: End-to-End Learning for Self-Driving Cars. Arxiv, 2016.

Next Time:
Direct Perception
(in two weeks)

Questions?